       Privacy Extensions for Stateless Address Autoconfiguration in IPv6

Status of This Memo

   This document specifies an Internet standards track protocol for the
   Internet community, and requests discussion and suggestions for
   improvements. Please refer to the current edition of the "Internet
   Official Protocol Standards" (STD 1) for the standardization state
   and status of this protocol. Distribution of this memo is unlimited.

Abstract

   Nodes use IPv6 stateless address autoconfiguration to generate
   addresses using a combination of locally available information and
   information advertised by routers. Addresses are formed by combining
   network prefixes with an interface identifier. On an interface that
   contains an embedded IEEE Identifier, the interface identifier is
   typically derived from it. On other interface types, the interface
   identifier is generated through other means, for example, via random
   number generation. This document describes an extension to IPv6
   stateless address autoconfiguration for interfaces whose interface
   identifier is derived from an IEEE identifier. Use of the extension
   causes nodes to generate global scope addresses from interface
   identifiers that change over time, even in cases where the interface
   contains an embedded IEEE identifier. Changing the interface
   identifier (and the global scope addresses generated from it) over
   time makes it more difficult for eavesdroppers and other information
   collectors to identify when different addresses used in different
   transactions actually correspond to the same node.

Table of Contents

1. Introduction

   Stateless address autoconfiguration [ADDRCONF] defines how an IPv6
   node generates addresses without the need for a Dynamic Host
   Configuration Protocol for IPv6 (DHCPv6) server. Some types of
   network interfaces come with an embedded IEEE Identifier (i.e., a
   link-layer MAC address), and in those cases, stateless address
   autoconfiguration uses the IEEE identifier to generate a 64-bit
   interface identifier [ADDRARCH]. By design, the interface identifier
   is likely to be globally unique when generated in this fashion. The
   interface identifier is in turn appended to a prefix to form a
   128-bit IPv6 address. Note that an IPv6 identifier does not
   necessarily have to be 64 bits in length, but the algorithm specified
   in this document is targeted towards 64-bit interface identifiers.

   All nodes combine interface identifiers (whether derived from an IEEE
   identifier or generated through some other technique) with the
   reserved link-local prefix to generate link-local addresses for their
   attached interfaces. Additional addresses can then be created by
   combining prefixes advertised in Router Advertisements via Neighbor
   Discovery [DISCOVERY] with the interface identifier.

   Not all nodes and interfaces contain IEEE identifiers. In such
   cases, an interface identifier is generated through some other means
   (e.g., at random), and the resultant interface identifier may not be
   globally unique and may also change over time. The focus of this
   document is on addresses derived from IEEE identifiers because
   tracking of individual devices, the concern being addressed here, is
   possible only in those cases where the interface identifier is
   globally unique and non-changing. The rest of this document assumes
   that IEEE identifiers are being used, but the techniques described
   may also apply to interfaces with other types of globally unique
   and/or persistent identifiers.

   This document discusses concerns associated with the embedding of
   non-changing interface identifiers within IPv6 addresses and
   describes extensions to stateless address autoconfiguration that can
   help mitigate those concerns for individual users and in environments
   where such concerns are significant. Section 2 provides background
   information on the issue. Section 3 describes a procedure for
   generating alternate interface identifiers and global scope
   addresses. Section 4 discusses implications of changing interface
   identifiers. The term "global scope addresses" is used in this
   document to collectively refer to "Global unicast addresses" as
   defined in [ADDRARCH] and "Unique local addresses" as defined in
   [ULA].

1.1. Conventions Used in This Document

   The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
   "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
   document are to be interpreted as described in [RFC2119].

1.2. Problem Statement

   Addresses generated using stateless address autoconfiguration
   [ADDRCONF] contain an embedded interface identifier, which remains
   constant over time. Anytime a fixed identifier is used in multiple
   contexts, it becomes possible to correlate seemingly unrelated
   activity using this identifier.

   The correlation can be performed by

   o An attacker who is in the path between the node in question and
      the peer(s) to which it is communicating, and who can view the
      IPv6 addresses present in the datagrams.

   o An attacker who can access the communication logs of the peers
      with which the node has communicated.

   Since the identifier is embedded within the IPv6 address, which is a
   fundamental requirement of communication, it cannot be easily hidden.
   This document proposes a solution to this issue by generating
   interface identifiers that vary over time.

   Note that an attacker, who is on path, may be able to perform
   significant correlation based on

   o The payload contents of the packets on the wire

   o The characteristics of the packets such as packet size and timing

   Use of temporary addresses will not prevent such payload-based
   correlation.

2. Background

   This section discusses the problem in more detail, provides context
   for evaluating the significance of the concerns in specific
   environments and makes comparisons with existing practices.

2.1. Extended Use of the Same Identifier

   The use of a non-changing interface identifier to form addresses is a
   specific instance of the more general case where a constant
   identifier is reused over an extended period of time and in multiple
   independent activities. Any time the same identifier is used in
   multiple contexts, it becomes possible for that identifier to be used
   to correlate seemingly unrelated activity. For example, a network
   sniffer placed strategically on a link across which all traffic to/
   from a particular host crosses could keep track of which destinations
   a node communicated with and at what times. Such information can in
   some cases be used to infer things, such as what hours an employee
   was active, when someone is at home, etc. Although it might appear
   that changing an address regularly in such environments would be
   desirable to lessen privacy concerns, it should be noted that the
   network prefix portion of an address also serves as a constant
   identifier. All nodes at, say, a home, would have the same network
   prefix, which identifies the topological location of those nodes.
   This has implications for privacy, though not at the same granularity
   as the concern that this document addresses. Specifically, all nodes
   within a home could be grouped together for the purposes of
   collecting information. If the network contains a very small number
   of nodes, say, just one, changing just the interface identifier will
   not enhance privacy at all, since the prefix serves as a constant
   identifier.

   One of the requirements for correlating seemingly unrelated
   activities is the use (and reuse) of an identifier that is
   recognizable over time within different contexts. IP addresses
   provide one obvious example, but there are more. Many nodes also
   have DNS names associated with their addresses, in which case the DNS
   name serves as a similar identifier. Although the DNS name
   associated with an address is more work to obtain (it may require a
   DNS query), the information is often readily available. In such
   cases, changing the address on a machine over time would do little to
   address the concerns raised in this document, unless the DNS name is
   changed as well (see Section 4).

   Web browsers and servers typically exchange "cookies" with each other
   [COOKIES]. Cookies allow Web servers to correlate a current activity
   with a previous activity. One common usage is to send back targeted
   advertising to a user by using the cookie supplied by the browser to

identify what earlier queries had been made (e.g., for what type of
information). Based on the earlier queries, advertisements can be
targeted to match the (assumed) interests of the end user.

The use of a constant identifier within an address is of special
concern because addresses are a fundamental requirement of
communication and cannot easily be hidden from eavesdroppers and
other parties. Even when higher layers encrypt their payloads,
addresses in packet headers appear in the clear. Consequently, if a
mobile host (e.g., laptop) accessed the network from several
different locations, an eavesdropper might be able to track the
movement of that mobile host from place to place, even if the upper
layer payloads were encrypted.

2.2. Address Usage in IPv4 Today

Addresses used in today's Internet are often non-changing in practice
for extended periods of time. In an increasing number of sites,
addresses are assigned statically and typically change infrequently.
Over the last few years, sites have begun moving away from static
allocation to dynamic allocation via DHCP [DHCP]. In theory, the
address a client gets via DHCP can change over time, but in practice
servers often return the same address to the same client (unless
addresses are in such short supply that they are reused immediately
by a different node when they become free). Thus, even within sites
using DHCP, clients frequently end up using the same address for
weeks to months at a time.

For home users accessing the Internet over dial-up lines, the
situation is generally different. Such users do not have permanent
connections and are often assigned temporary addresses each time they
connect to their ISP. Consequently, the addresses they use change
frequently over time and are shared among a number of different
users. Thus, an address does not reliably identify a particular
device over time spans of more than a few minutes.

A more interesting case concerns always-on connections (e.g., cable
modems, ISDN, DSL, etc.) that result in a home site using the same
address for extended periods of time. This is a scenario that is
just starting to become common in IPv4 and promises to become more of
a concern as always-on Internet connectivity becomes widely
available.

Finally, it should be noted that nodes that need a (non-changing) DNS
name generally have static addresses assigned to them to simplify the
configuration of DNS servers. Although Dynamic DNS [DDNS] can be
used to update the DNS dynamically, it may not always be available
depending on the administrative policy. In addition, changing an

      address but keeping the same DNS name does not really address the
      underlying concern, since the DNS name becomes a non-changing
      identifier. Servers generally require a DNS name (so clients can
      connect to them), and clients often do as well (e.g., some servers
      refuse to speak to a client whose address cannot be mapped into a DNS
      name that also maps back into the same address). Section 4 describes
      one approach to this issue.

2.3. The Concern with IPv6 Addresses

      The division of IPv6 addresses into distinct topology and interface
      identifier portions raises an issue new to IPv6 in that a fixed
      portion of an IPv6 address (i.e., the interface identifier) can
      contain an identifier that remains constant even when the topology
      portion of an address changes (e.g., as the result of connecting to a
      different part of the Internet). In IPv4, when an address changes,
      the entire address (including the local part of the address) usually
      changes. It is this new issue that this document addresses.

      If addresses are generated from an interface identifier, a home
      user's address could contain an interface identifier that remains the
      same from one dial-up session to the next, even if the rest of the
      address changes. The way PPP is used today, however, PPP servers
      typically unilaterally inform the client what address they are to use
      (i.e., the client doesn't generate one on its own). This practice,
      if continued in IPv6, would avoid the concerns that are the focus of
      this document.

      A more troubling case concerns mobile devices (e.g., laptops, PDAs,
      etc.) that move topologically within the Internet. Whenever they
      move, they form new addresses for their current topological point of
      attachment. This is typified today by the "road warrior" who has
      Internet connectivity both at home and at the office. While the
      node's address changes as it moves, the interface identifier
      contained within the address remains the same (when derived from an
      IEEE Identifier). In such cases, the interface identifier can be
      used to track the movement and usage of a particular machine. For
      example, a server that logs usage information together with source
      addresses, is also recording the interface identifier since it is
      embedded within an address. Consequently, any data-mining technique
      that correlates activity based on addresses could easily be extended
      to do the same using the interface identifier. This is of particular
      concern with the expected proliferation of next-generation network-
      connected devices (e.g., PDAs, cell phones, etc.) in which large
      numbers of devices are, in practice, associated with individual users
      (i.e., not shared). Thus, the interface identifier embedded within
      an address could be used to track activities of an individual, even
      as they move topologically within the Internet.

   In summary, IPv6 addresses on a given interface generated via
   Stateless Autoconfiguration contain the same interface identifier,
   regardless of where within the Internet the device connects. This
   facilitates the tracking of individual devices (and thus,
   potentially, users). The purpose of this document is to define
   mechanisms that eliminate this issue in those situations where it is
   a concern.

2.4. Possible Approaches

   One way to avoid having a static non-changing address is to use
   DHCPv6 [DHCPV6] for obtaining addresses. Section 12 of [DHCPV6]
   discusses the use of DHCPv6 for the assignment and management of
   "temporary addresses", which are never renewed and provide the same
   property of temporary addresses described in this document with
   regards to the privacy concern.

   Another approach, compatible with the stateless address
   autoconfiguration architecture, would be to change the interface
   identifier portion of an address over time and generate new addresses
   from the interface identifier for some address scopes. Changing the
   interface identifier can make it more difficult to look at the IP
   addresses in independent transactions and identify which ones
   actually correspond to the same node, both in the case where the
   routing prefix portion of an address changes and when it does not.

   Many machines function as both clients and servers. In such cases,
   the machine would need a DNS name for its use as a server. Whether
   the address stays fixed or changes has little privacy implication
   since the DNS name remains constant and serves as a constant
   identifier. When acting as a client (e.g., initiating
   communication), however, such a machine may want to vary the
   addresses it uses. In such environments, one may need multiple
   addresses: a "public" (i.e., non-secret) server address, registered
   in the DNS, that is used to accept incoming connection requests from
   other machines, and a "temporary" address used to shield the identity
   of the client when it initiates communication. These two cases are
   roughly analogous to telephone numbers and caller ID, where a user
   may list their telephone number in the public phone book, but disable
   the display of its number via caller ID when initiating calls.

   To make it difficult to make educated guesses as to whether two
   different interface identifiers belong to the same node, the
   algorithm for generating alternate identifiers must include input
   that has an unpredictable component from the perspective of the
   outside entities that are collecting information. Picking
   identifiers from a pseudo-random sequence suffices, so long as the
   specific sequence cannot be determined by an outsider examining

information that is readily available or easily determinable (e.g.,
by examining packet contents). This document proposes the generation
of a pseudo-random sequence of interface identifiers via an MD5 hash.
Periodically, the next interface identifier in the sequence is
generated, a new set of temporary addresses is created, and the
previous temporary addresses are deprecated to discourage their
further use. The precise pseudo-random sequence depends on both a
random component and the globally unique interface identifier (when
available), to increase the likelihood that different nodes generate
different sequences.

3. Protocol Description

   The goal of this section is to define procedures that:

   1. Do not result in any changes to the basic behavior of addresses
       generated via stateless address autoconfiguration [ADDRCONF].

   2. Create additional addresses based on a random interface
       identifier for the purpose of initiating outgoing sessions.
       These "random" or temporary addresses would be used for a short
       period of time (hours to days) and would then be deprecated.
       Deprecated address can continue to be used for already
       established connections, but are not used to initiate new
       connections. New temporary addresses are generated periodically
       to replace temporary addresses that expire, with the exact time
       between address generation a matter of local policy.

   3. Produce a sequence of temporary global scope addresses from a
       sequence of interface identifiers that appear to be random in the
       sense that it is difficult for an outside observer to predict a
       future address (or identifier) based on a current one, and it is
       difficult to determine previous addresses (or identifiers)
       knowing only the present one.

   4. By default, generate a set of addresses from the same
       (randomized) interface identifier, one address for each prefix
       for which a global address has been generated via stateless
       address autoconfiguration. Using the same interface identifier
       to generate a set of temporary addresses reduces the number of IP
       multicast groups a host must join. Nodes join the solicited-node
       multicast address for each unicast address they support, and
       solicited-node addresses are dependent only on the low-order bits
       of the corresponding address. This default behavior was made to
       address the concern that a node that joins a large number of
       multicast groups may be required to put its interface into
       promiscuous mode, resulting in possible reduced performance.

A node highly concerned about privacy MAY use different interface identifiers on different prefixes, resulting in a set of global addresses that cannot be easily tied to each other. For example a node MAY create different interface identifiers I1, I2, and I3 for use with different prefixes P1, P2, and P3 on the same interface.

3.1. Assumptions

The following algorithm assumes that each interface maintains an associated randomized interface identifier. When temporary addresses are generated, the current value of the associated randomized interface identifier is used. While the same identifier can be used to create more than one temporary address, the value SHOULD change over time as described in Section 3.5.

The algorithm also assumes that, for a given temporary address, an implementation can determine the prefix from which it was generated. When a temporary address is deprecated, a new temporary address is generated. The specific valid and preferred lifetimes for the new address are dependent on the corresponding lifetime values set for the prefix from which it was generated.

Finally, this document assumes that when a node initiates outgoing communication, temporary addresses can be given preference over public addresses when the device is configured to do so. [ADDR_SELECT] mandates implementations to provide a mechanism, which allows an application to configure its preference for temporary addresses over public addresses. It also allows for an implementation to prefer temporary addresses by default, so that the connections initiated by the node can use temporary addresses without requiring application-specific enablement. This document also assumes that an API will exist that allows individual applications to indicate whether they prefer to use temporary or public addresses and override the system defaults.

3.2. Generation of Randomized Interface Identifiers

We describe two approaches for the generation and maintenance of the randomized interface identifier. The first assumes the presence of stable storage that can be used to record state history for use as input into the next iteration of the algorithm across system restarts. A second approach addresses the case where stable storage is unavailable and there is a need to generate randomized interface identifiers without previous state.

The random interface identifier generation algorithm, as described in
this document, uses MD5 as the hash algorithm. The node MAY use
another algorithm instead of MD5 to produce the random interface
identifier.

3.2.1. When Stable Storage Is Present

The following algorithm assumes the presence of a 64-bit "history
value" that is used as input in generating a randomized interface
identifier. The very first time the system boots (i.e., out-of-the-
box), a random value SHOULD be generated using techniques that help
ensure the initial value is hard to guess [RANDOM]. Whenever a new
interface identifier is generated, a value generated by the
computation is saved in the history value for the next iteration of
the algorithm.

A randomized interface identifier is created as follows:

1.   Take the history value from the previous iteration of this
     algorithm (or a random value if there is no previous value) and
     append to it the interface identifier generated as described in
     [ADDRARCH].

2.   Compute the MD5 message digest [MD5] over the quantity created in
     the previous step.

3.   Take the leftmost 64-bits of the MD5 digest and set bit 6 (the
     leftmost bit is numbered 0) to zero. This creates an interface
     identifier with the universal/local bit indicating local
     significance only.

4.   Compare the generated identifier against a list of reserved
     interface identifiers and to those already assigned to an address
     on the local device. In the event that an unacceptable
     identifier has been generated, the node MUST restart the process
     at step 1 above, using the rightmost 64 bits of the MD5 digest
     obtained in step 2 in place of the history value in step 1.

5.   Save the generated identifier as the associated randomized
     interface identifier.

6.   Take the rightmost 64-bits of the MD5 digest computed in step 2)
     and save them in stable storage as the history value to be used
     in the next iteration of the algorithm.

MD5 was chosen for convenience, and because its particular properties
were adequate to produce the desired level of randomization. The
node MAY use another algorithm instead of MD5 to produce the random
interface identifier

In theory, generating successive randomized interface identifiers
using a history scheme as above has no advantages over generating
them at random. In practice, however, generating truly random
numbers can be tricky. Use of a history value is intended to avoid
the particular scenario where two nodes generate the same randomized
interface identifier, both detect the situation via DAD, but then
proceed to generate identical randomized interface identifiers via
the same (flawed) random number generation algorithm. The above
algorithm avoids this problem by having the interface identifier
(which will often be globally unique) used in the calculation that
generates subsequent randomized interface identifiers. Thus, if two
nodes happen to generate the same randomized interface identifier,
they should generate different ones on the follow-up attempt.

3.2.2. In The Absence of Stable Storage

In the absence of stable storage, no history value will be available
across system restarts to generate a pseudo-random sequence of
interface identifiers. Consequently, the initial history value used
above SHOULD be generated at random. A number of techniques might be
appropriate. Consult [RANDOM] for suggestions on good sources for
obtaining random numbers. Note that even though machines may not
have stable storage for storing a history value, they will in many
cases have configuration information that differs from one machine to
another (e.g., user identity, security keys, serial numbers, etc.).
One approach to generating a random initial history value in such
cases is to use the configuration information to generate some data
bits (which may remain constant for the life of the machine, but will
vary from one machine to another), append some random data, and
compute the MD5 digest as before.

3.2.3. Alternate Approaches

Note that there are other approaches to generate random interface
identifiers, albeit with different goals and applicability. One such
approach is Cryptographically Generated Addresses (CGAs) [CGA], which
generate a random interface identifier based on the public key of the
node. The goal of CGAs is to prove ownership of an address and to
prevent spoofing and stealing of existing IPv6 addresses. They are
used for securing neighbor discovery using [SEND]. The CGA random
interface identifier generation algorithm may not be suitable for
privacy addresses because of the following properties:

   o It requires the node to have a public key. This means that the
      node can still be identified by its public key.

   o The random interface identifier process is computationally
      intensive and hence discourages frequent regeneration.

3.3. Generating Temporary Addresses

   [ADDRCONF] describes the steps for generating a link-local address
   when an interface becomes enabled as well as the steps for generating
   addresses for other scopes. This document extends [ADDRCONF] as
   follows. When processing a Router Advertisement with a Prefix
   Information option carrying a global scope prefix for the purposes of
   address autoconfiguration (i.e., the A bit is set), the node MUST
   perform the following steps:

   1. Process the Prefix Information Option as defined in [ADDRCONF],
       either creating a new public address or adjusting the lifetimes
       of existing addresses, both public and temporary. If a received
       option will extend the lifetime of a public address, the
       lifetimes of temporary addresses should be extended, subject to
       the overall constraint that no temporary addresses should ever
       remain "valid" or "preferred" for a time longer than
       (TEMP_VALID_LIFETIME) or (TEMP_PREFERRED_LIFETIME -
       DESYNC_FACTOR), respectively. The configuration variables
       TEMP_VALID_LIFETIME and TEMP_PREFERRED_LIFETIME correspond to
       approximate target lifetimes for temporary addresses.

   2. One way an implementation can satisfy the above constraints is to
       associate with each temporary address a creation time (called
       CREATION_TIME) that indicates the time at which the address was
       created. When updating the preferred lifetime of an existing
       temporary address, it would be set to expire at whichever time is
       earlier: the time indicated by the received lifetime or
       (CREATION_TIME + TEMP_PREFERRED_LIFETIME - DESYNC_FACTOR). A
       similar approach can be used with the valid lifetime.

   3. When a new public address is created as described in [ADDRCONF],
       the node SHOULD also create a new temporary address.

   4. When creating a temporary address, the lifetime values MUST be
       derived from the corresponding prefix as follows:

       * Its Valid Lifetime is the lower of the Valid Lifetime of the
          public address or TEMP_VALID_LIFETIME.

   * Its Preferred Lifetime is the lower of the Preferred Lifetime
      of the public address or TEMP_PREFERRED_LIFETIME -
      DESYNC_FACTOR.

   5. A temporary address is created only if this calculated Preferred
       Lifetime is greater than REGEN_ADVANCE time units. In
       particular, an implementation MUST NOT create a temporary address
       with a zero Preferred Lifetime.

   6. New temporary addresses MUST be created by appending the
       interface's current randomized interface identifier to the prefix
       that was received.

   7. The node MUST perform duplicate address detection (DAD) on the
       generated temporary address. If DAD indicates the address is
       already in use, the node MUST generate a new randomized interface
       identifier as described in Section 3.2 above, and repeat the
       previous steps as appropriate up to TEMP_IDGEN_RETRIES times. If
       after TEMP_IDGEN_RETRIES consecutive attempts no non-unique
       address was generated, the node MUST log a system error and MUST
       NOT attempt to generate temporary addresses for that interface.
       Note that DAD MUST be performed on every unicast address
       generated from this randomized interface identifier.

3.4. Expiration of Temporary Addresses

   When a temporary address becomes deprecated, a new one MUST be
   generated. This is done by repeating the actions described in
   Section 3.3, starting at step 3). Note that, except for the
   transient period when a temporary address is being regenerated, in
   normal operation at most one temporary address per prefix should be
   in a non-deprecated state at any given time on a given interface.
   Note that if a temporary address becomes deprecated as result of
   processing a Prefix Information Option with a zero Preferred
   Lifetime, then a new temporary address MUST NOT be generated. To
   ensure that a preferred temporary address is always available, a new
   temporary address SHOULD be regenerated slightly before its
   predecessor is deprecated. This is to allow sufficient time to avoid
   race conditions in the case where generating a new temporary address
   is not instantaneous, such as when duplicate address detection must
   be run. The node SHOULD start the address regeneration process
   REGEN_ADVANCE time units before a temporary address would actually be
   deprecated.

   As an optional optimization, an implementation MAY remove a
   deprecated temporary address that is not in use by applications or
   upper layers as detailed in Section 6.

3.5. Regeneration of Randomized Interface Identifiers

   The frequency at which temporary addresses changes depends on how a
   device is being used (e.g., how frequently it initiates new
   communication) and the concerns of the end user. The most egregious
   privacy concerns appear to involve addresses used for long periods of
   time (weeks to months to years). The more frequently an address
   changes, the less feasible collecting or coordinating information
   keyed on interface identifiers becomes. Moreover, the cost of
   collecting information and attempting to correlate it based on
   interface identifiers will only be justified if enough addresses
   contain non-changing identifiers to make it worthwhile. Thus, having
   large numbers of clients change their address on a daily or weekly
   basis is likely to be sufficient to alleviate most privacy concerns.

   There are also client costs associated with having a large number of
   addresses associated with a node (e.g., in doing address lookups, the
   need to join many multicast groups, etc.). Thus, changing addresses
   frequently (e.g., every few minutes) may have performance
   implications.

   Nodes following this specification SHOULD generate new temporary
   addresses on a periodic basis. This can be achieved automatically by
   generating a new randomized interface identifier at least once every
   (TEMP_PREFERRED_LIFETIME - REGEN_ADVANCE - DESYNC_FACTOR) time units.
   As described above, generating a new temporary address REGEN_ADVANCE
   time units before a temporary address becomes deprecated produces
   addresses with a preferred lifetime no larger than
   TEMP_PREFERRED_LIFETIME. The value DESYNC_FACTOR is a random value
   (different for each client) that ensures that clients don't
   synchronize with each other and generate new addresses at exactly the
   same time. When the preferred lifetime expires, a new temporary
   address MUST be generated using the new randomized interface
   identifier.

   Because the precise frequency at which it is appropriate to generate
   new addresses varies from one environment to another, implementations
   SHOULD provide end users with the ability to change the frequency at
   which addresses are regenerated. The default value is given in
   TEMP_PREFERRED_LIFETIME and is one day. In addition, the exact time
   at which to invalidate a temporary address depends on how
   applications are used by end users. Thus, the suggested default
   value of one week (TEMP_VALID_LIFETIME) may not be appropriate in all
   environments. Implementations SHOULD provide end users with the
   ability to override both of these default values.

Finally, when an interface connects to a new link, a new randomized interface identifier SHOULD be generated immediately together with a new set of temporary addresses. If a device moves from one ethernet to another, generating a new set of temporary addresses from a different randomized interface identifier ensures that the device uses different randomized interface identifiers for the temporary addresses associated with the two links, making it more difficult to correlate addresses from the two different links as being from the same node. The node MAY follow any process available to it, to determine that the link change has occurred. One such process is described by Detecting Network Attachment [DNA].

3.6. Deployment Considerations

Devices implementing this specification MUST provide a way for the end user to explicitly enable or disable the use of temporary addresses. In addition, a site might wish to disable the use of temporary addresses in order to simplify network debugging and operations. Consequently, implementations SHOULD provide a way for trusted system administrators to enable or disable the use of temporary addresses.

Additionally, sites might wish to selectively enable or disable the use of temporary addresses for some prefixes. For example, a site might wish to disable temporary address generation for "Unique local" [ULA] prefixes while still generating temporary addresses for all other global prefixes. Another site might wish to enable temporary address generation only for the prefixes 2001::/16 and 2002::/16, while disabling it for all other prefixes. To support this behavior, implementations SHOULD provide a way to enable and disable generation of temporary addresses for specific prefix subranges. This per-prefix setting SHOULD override the global settings on the node with respect to the specified prefix subranges. Note that the pre-prefix setting can be applied at any granularity, and not necessarily on a per-subnet basis.

The use of temporary addresses may cause unexpected difficulties with some applications. As described below, some servers refuse to accept communications from clients for which they cannot map the IP address into a DNS name. In addition, some applications may not behave robustly if temporary addresses are used and an address expires before the application has terminated, or if it opens multiple sessions, but expects them to all use the same addresses. Consequently, the use of temporary addresses SHOULD be disabled by default in order to minimize potential disruptions. Individual applications, which have specific knowledge about the normal duration of connections, MAY override this as appropriate.

   If a very small number of nodes (say, only one) use a given prefix
   for extended periods of time, just changing the interface identifier
   part of the address may not be sufficient to ensure privacy, since
   the prefix acts as a constant identifier. The procedures described
   in this document are most effective when the prefix is reasonably non
   static or is used by a fairly large number of nodes.

4. Implications of Changing Interface Identifiers

   The IPv6 addressing architecture goes to some lengths to ensure that
   interface identifiers are likely to be globally unique where easy to
   do so. The widespread use of temporary addresses may result in a
   significant fraction of Internet traffic not using addresses in which
   the interface identifier portion is globally unique. Consequently,
   usage of the algorithms in this document may complicate providing
   such a future flexibility, if global uniqueness is necessary.

   The desires of protecting individual privacy versus the desire to
   effectively maintain and debug a network can conflict with each
   other. Having clients use addresses that change over time will make
   it more difficult to track down and isolate operational problems.
   For example, when looking at packet traces, it could become more
   difficult to determine whether one is seeing behavior caused by a
   single errant machine, or by a number of them.

   Some servers refuse to grant access to clients for which no DNS name
   exists. That is, they perform a DNS PTR query to determine the DNS
   name, and may then also perform an AAAA query on the returned name to
   verify that the returned DNS name maps back into the address being
   used. Consequently, clients not properly registered in the DNS may
   be unable to access some services. As noted earlier, however, a
   node's DNS name (if non-changing) serves as a constant identifier.
   The wide deployment of the extension described in this document could
   challenge the practice of inverse-DNS-based "authentication," which
   has little validity, though it is widely implemented. In order to
   meet server challenges, nodes could register temporary addresses in
   the DNS using random names (for example, a string version of the
   random address itself).

   Use of the extensions defined in this document may complicate
   debugging and other operational troubleshooting activities.
   Consequently, it may be site policy that temporary addresses should
   not be used. Consequently, implementations MUST provide a method for
   the end user or trusted administrator to override the use of
   temporary addresses.

5. Defined Constants

    Constants defined in this document include:

    TEMP_VALID_LIFETIME -- Default value: 1 week. Users should be able
    to override the default value.

    TEMP_PREFERRED_LIFETIME -- Default value: 1 day. Users should be
    able to override the default value.

    REGEN_ADVANCE -- 5 seconds

    MAX_DESYNC_FACTOR -- 10 minutes. Upper bound on DESYNC_FACTOR.

    DESYNC_FACTOR -- A random value within the range 0 -
    MAX_DESYNC_FACTOR. It is computed once at system start (rather than
    each time it is used) and must never be greater than
    (TEMP_VALID_LIFETIME - REGEN_ADVANCE).

    TEMP_IDGEN_RETRIES -- Default value: 3

6. Future Work

    An implementation might want to keep track of which addresses are
    being used by upper layers so as to be able to remove a deprecated
    temporary address from internal data structures once no upper layer
    protocols are using it (but not before). This is in contrast to
    current approaches where addresses are removed from an interface when
    they become invalid [ADDRCONF], independent of whether or not upper
    layer protocols are still using them. For TCP connections, such
    information is available in control blocks. For UDP-based
    applications, it may be the case that only the applications have
    knowledge about what addresses are actually in use. Consequently, an
    implementation generally will need to use heuristics in deciding when
    an address is no longer in use.

    The determination as to whether to use public versus temporary
    addresses can in some cases only be made by an application. For
    example, some applications may always want to use temporary
    addresses, while others may want to use them only in some
    circumstances or not at all. Suitable API extensions will likely
    need to be developed to enable individual applications to indicate
    with sufficient granularity their needs with regards to the use of
    temporary addresses. Recommendations on DNS practices to avoid the
    problem described in Section 4 when reverse DNS lookups fail may be
    needed. [DNSOP] contains a more detailed discussion of the DNS-
    related issues.

   While this document discusses ways of obscuring a user's permanent IP
   address, the method described is believed to be ineffective against
   sophisticated forms of traffic analysis. To increase effectiveness,
   one may need to consider use of more advanced techniques, such as
   Onion Routing [ONION].

7. Security Considerations

   Ingress filtering has been and is being deployed as a means of
   preventing the use of spoofed source addresses in Distributed Denial
   of Service (DDoS) attacks. In a network with a large number of
   nodes, new temporary addresses are created at a fairly high rate.
   This might make it difficult for ingress filtering mechanisms to
   distinguish between legitimately changing temporary addresses and
   spoofed source addresses, which are "in-prefix" (using a
   topologically correct prefix and non-existent interface ID). This
   can be addressed by using access control mechanisms on a per-address
   basis on the network egress point.

8. Significant Changes from RFC 3041

   This section summarizes the changes in this document relative to RFC
   3041 that an implementer of RFC 3041 should be aware of.

   1. Excluded certain interface identifiers from the range of
       acceptable interface identifiers. Interface IDs such as those
       for reserved anycast addresses [RFC2526], etc.

   2. Added a configuration knob that provides the end user with a way
       to enable or disable the use of temporary addresses on a per-
       prefix basis.

   3. Added a check for denial of service attacks using low valid
       lifetimes in router advertisements.

   4. DAD is now run on all temporary addresses, not just the first one
       generated from an interface identifier.

   5. Changed the default setting for usage of temporary addresses to
       be disabled.

   6. The node is now allowed to generate different interface
       identifiers for different prefixes, if it so desires.

   7. The algorithm used for generating random interface identifiers is
       no longer restricted to just MD5.

   8. Reduced default number of retries to 3 and added a configuration
       variable.

   9. Router advertisement (RA) processing algorithm is no longer
       included in the document, and is replaced by a reference to
       [ADDRCONF].

9. Acknowledgments

   Rich Draves and Thomas Narten were the authors of RFC 3041. They
   would like to acknowledge the contributions of the ipv6 working group
   and, in particular, Ran Atkinson, Matt Crawford, Steve Deering,
   Allison Mankin, and Peter Bieringer.

   Suresh Krishnan was the sole author of this version of the document.
   He would like to acknowledge the contributions of the ipv6 working
   group and, in particular, Jari Arkko, Pekka Nikander, Pekka Savola,
   Francis Dupont, Brian Haberman, Tatuya Jinmei, and Margaret Wasserman
   for their detailed comments.

10. References

10.1. Normative References

   [ADDRARCH]   Hinden, R. and S. Deering, "IP Version 6 Addressing
                  Architecture", RFC 4291, February 2006.

   [ADDRCONF]   Thomson, S., Narten, T., and T. Jinmei, "IPv6
                  Stateless Address Autoconfiguration", RFC 4862,
                  September 2007.

   [DISCOVERY]   Narten, T., Nordmark, E., Simpson, W., and H. Soliman,
                  "Neighbor Discovery for IP version 6 (IPv6)",
                  RFC 4861, September 2007.

   [MD5]   Rivest, R., "The MD5 Message-Digest Algorithm",
                  RFC 1321, April 1992.

   [RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate
                  Requirement Levels", RFC 2119, March 1997.

10.2. Informative References

   [ADDR_SELECT] Draves, R., "Default Address Selection for Internet
                  Protocol version 6 (IPv6)", RFC 3484, February 2003.

   [CGA]   Aura, T., "Cryptographically Generated Addresses
                  (CGA)", RFC 3972, March 2005.

[COOKIES]   Kristol, D. and L. Montulli, "HTTP State Management
              Mechanism", RFC 2965, October 2000.

[DDNS]   Vixie, P., Thomson, S., Rekhter, Y., and J. Bound,
              "Dynamic Updates in the Domain Name System (DNS
              UPDATE)", RFC 2136, April 1997.

[DHCP]   Droms, R., "Dynamic Host Configuration Protocol",
              RFC 2131, March 1997.

[DHCPV6]   Droms, R., Bound, J., Volz, B., Lemon, T., Perkins,
              C., and M. Carney, "Dynamic Host Configuration
              Protocol for IPv6 (DHCPv6)", RFC 3315, July 2003.

[DNA]   Choi, JH. and G. Daley, "Goals of Detecting Network
              Attachment in IPv6", RFC 4135, August 2005.

[DNSOP]   Durand, A., Ihren, J., and P. Savola, "Operational
              Considerations and Issues with IPv6 DNS", RFC 4472,
              April 2006.

[ONION]   Reed, MGR., Syverson, PFS., and DMG. Goldschlag,
              "Proxies for Anonymous Routing", Proceedings of the
              12th Annual Computer Security Applications Conference,
              San Diego, CA, December 1996.

[RANDOM]   Eastlake, D., Schiller, J., and S. Crocker,
              "Randomness Requirements for Security", BCP 106,
              RFC 4086, June 2005.

[RFC2526]   Johnson, D. and S. Deering, "Reserved IPv6 Subnet
              Anycast Addresses", RFC 2526, March 1999.

[SEND]   Arkko, J., Kempf, J., Zill, B., and P. Nikander,
              "SEcure Neighbor Discovery (SEND)", RFC 3971,
              March 2005.

[ULA]   Hinden, R. and B. Haberman, "Unique Local IPv6 Unicast
              Addresses", RFC 4193, October 2005.

Authors' Addresses

   Thomas Narten
   IBM Corporation
   P.O. Box 12195
   Research Triangle Park, NC
   USA

   EMail: narten@us.ibm.com

   Richard Draves
   Microsoft Research
   One Microsoft Way
   Redmond, WA
   USA

   EMail: richdr@microsoft.com

   Suresh Krishnan
   Ericsson Research
   8400 Decarie Blvd.
   Town of Mount Royal, QC
   Canada

   EMail: suresh.krishnan@ericsson.com