          A Per-Domain Path Computation Method for Establishing Inter-Domain
               Traffic Engineering (TE) Label Switched Paths (LSPs)

Status of This Memo

   This document specifies an Internet standards track protocol for the
   Internet community, and requests discussion and suggestions for
   improvements. Please refer to the current edition of the "Internet
   Official Protocol Standards" (STD 1) for the standardization state
   and status of this protocol. Distribution of this memo is unlimited.

Abstract

   This document specifies a per-domain path computation technique for
   establishing inter-domain Traffic Engineering (TE) Multiprotocol
   Label Switching (MPLS) and Generalized MPLS (GMPLS) Label Switched
   Paths (LSPs). In this document, a domain refers to a collection of
   network elements within a common sphere of address management or path
   computational responsibility such as Interior Gateway Protocol (IGP)
   areas and Autonomous Systems.

   Per-domain computation applies where the full path of an inter-domain
   TE LSP cannot be or is not determined at the ingress node of the TE
   LSP, and is not signaled across domain boundaries. This is most
   likely to arise owing to TE visibility limitations. The signaling
   message indicates the destination and nodes up to the next domain
   boundary. It may also indicate further domain boundaries or domain
   identifiers. The path through each domain, possibly including the
   choice of exit point from the domain, must be determined within the
   domain.

Table of Contents

1. Introduction

   The requirements for inter-domain Traffic Engineering (inter-area and
   inter-AS TE) have been developed by the Traffic Engineering Working
   Group and have been stated in [RFC4105] and [RFC4216]. The framework
   for inter-domain MPLS Traffic Engineering has been provided in
   [RFC4726].

   Some of the mechanisms used to establish and maintain inter-domain TE
   LSPs are specified in [RFC5151] and [RFC5150].

   This document exclusively focuses on the path computation aspects and
   defines a method for establishing inter-domain TE LSPs where each
   node in charge of computing a section of an inter-domain TE LSP path
   is always along the path of such a TE LSP.

   When the visibility of an end-to-end complete path spanning multiple
   domains is not available at the Head-end LSR (the LSR that initiated
   the TE LSP), one approach described in this document consists of
   using a per-domain path computation technique during LSP setup to
   determine the inter-domain TE LSP as it traverses multiple domains.

The mechanisms proposed in this document are also applicable to MPLS TE domains other than IGP areas and ASs.

The solution described in this document does not attempt to address all the requirements specified in [RFC4105] and [RFC4216]. This is acceptable according to [RFC4216], which indicates that a solution may be developed to address a particular deployment scenario and might, therefore, not meet all requirements for other deployment scenarios.

It must be pointed out that the inter-domain path computation technique proposed in this document is one among many others. The choice of the appropriate technique must be driven by the set of requirements for the path attributes and the applicability to a particular technique with respect to the deployment scenario. For example, if the requirement is to get an end-to-end constraint-based shortest path across multiple domains, then a mechanism using one or more distributed PCEs could be used to compute the shortest path across different domains (see [RFC4655]). Other off-line mechanisms for path computation are not precluded either. Note also that a Service Provider may elect to use different inter-domain path computation techniques for different TE LSP types.

2. Terminology

Terminology used in this document:

AS: Autonomous System.

ABR: Area Border Router, a router used to connect two IGP areas (areas in OSPF or levels in Intermediate System to Intermediate System (IS-IS)).

ASBR: Autonomous System Border Router, a router used to connect together ASs of a different or the same Service Provider via one or more inter-AS links.

Boundary LSR: A boundary LSR is either an ABR in the context of inter-area TE or an ASBR in the context of inter-AS TE.

ERO: Explicit Route Object.

IGP: Interior Gateway Protocol.

Inter-AS TE LSP: A TE LSP that crosses an AS boundary.

Inter-area TE LSP: A TE LSP that crosses an IGP area.

   LSR: Label Switching Router.

   LSP: Label Switched Path.

   TE LSP: Traffic Engineering Label Switched Path.

   PCE: Path Computation Element, an entity (component, application, or
   network node) that is capable of computing a network path or route
   based on a network graph and applying computational constraints.

   TED: Traffic Engineering Database.

   The notion of contiguous, stitched, and nested TE LSPs is defined in
   [RFC4726] and will not be repeated here.

2.1. Requirements Language

   The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
   "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
   document are to be interpreted as described in RFC 2119 [RFC2119].

3. General Assumptions

3.1. Common Assumptions

   - Each domain in all the examples below is assumed to be capable of
     doing Traffic Engineering (i.e., running OSPF-TE or ISIS-TE and
     RSVP-TE (Resource Reservation Protocol Traffic Engineering)). A
     domain may itself comprise multiple other domains, e.g., an AS may
     itself be composed of several other sub-ASs (BGP confederations) or
     areas/levels. In this case, the path computation technique
     described for inter-area and inter-AS MPLS Traffic Engineering
     applies recursively.

   - The inter-domain TE LSPs are signaled using RSVP-TE ([RFC3209] and
     [RFC3473]).

   - The path (specified by an ERO (Explicit Route Object) in an RSVP-TE
     Path message) for an inter-domain TE LSP may be signaled as a set
     of (loose and/or strict) hops.

   - The hops may identify:

     * The complete strict path end-to-end across different domains

     * The complete strict path in the source domain followed by
        boundary LSRs (or domain identifiers, e.g., AS numbers)

      * The complete list of boundary LSRs along the path

      * The current boundary LSR and the LSP destination

   The set of (loose or strict) hops can be either statically configured
   on the Head-end LSR or dynamically computed. A per-domain path
   computation method is defined in this document with an optional
   auto-discovery mechanism (e.g., based on IGP, BGP, policy routing
   information) yielding the next-hop boundary node (domain exit point,
   such as an Area Border Router (ABR) or an Autonomous System Border
   Router (ASBR)) along the path as the TE LSP is being signaled, along
   with potential crankback mechanisms. Alternatively, the domain exit
   points may be statically configured on the Head-end LSR, in which
   case next-hop boundary node auto-discovery would not be required.

   - Boundary LSRs are assumed to be capable of performing local path
     computation for expansion of a loose next hop in the signaled ERO
     if the path is not signaled by the Head-end LSR as a set of strict
     hops or if the strict hop is an abstract node (e.g., an AS). In
     any case, no topology or resource information needs to be
     distributed between domains (as mandated per [RFC4105] and
     [RFC4216]), which is critical to preserve IGP/BGP scalability and
     confidentiality in the case of TE LSPs spanning multiple routing
     domains.

   - The paths for the intra-domain Hierarchical LSP (H-LSP) or Stitched
     LSP (S-LSP) or for a contiguous TE LSP within the domain may be
     pre-configured or computed dynamically based on the arriving
     inter-domain LSP setup request (depending on the requirements of
     the transit domain). Note that this capability is explicitly
     specified as a requirement in [RFC4216]. When the paths for the
     H-LSP/S-LSP are pre-configured, the constraints as well as other
     parameters like a local protection scheme for the intra-domain H-
     LSP/S-LSP are also pre-configured.

   - While certain constraints like bandwidth can be used across
     different domains, certain other TE constraints like resource
     affinity, color, metric, etc. as listed in [RFC2702] may need to be
     translated at domain boundaries. If required, it is assumed that,
     at the domain boundary LSRs, there will exist some sort of local
     mapping based on policy agreement in order to translate such
     constraints across domain boundaries. It is expected that such an
     assumption particularly applies to inter-AS TE: for example, the
     local mapping would be similar to the inter-AS TE agreement
     enforcement polices stated in [RFC4216].

      - The procedures defined in this document are applicable to any node
        (not just a boundary node) that receives a Path message with an ERO
        that constrains a loose hop or an abstract node that is not a
        simple abstract node (that is, an abstract node that identifies
        more than one LSR).

3.2. Example of Topology for the Inter-Area TE Case

   The following example will be used for the inter-area TE case in this
   document.

```
              <-area 1-><-- area 0 --><--- area 2 --->
              ------ABR1------------ABR3-------
              |    /    |              |  \    |
              R0--X1    |              |   X2---X3--R1
              |         |              | /      |
              ------ABR2----------ABR4-------
              <=========== Inter-area TE LSP =======>
```

        Figure 1 - Example of topology for the inter-area TE case

   Description of Figure 1:

   - ABR1, ABR2, ABR3, and ABR4 are ABRs.
   - X1 is an LSR in area 1.
   - X2 and X3 are LSRs in area 2.
   - An inter-area TE LSP T0 originated at R0 in area 1 and terminated
     at R1 in area 2.

   Notes:

   - The terminology used in the example above corresponds to OSPF, but
     the path computation technique proposed in this document equally
     applies to the case of an IS-IS multi-level network.

   - Just a few routers in each area are depicted in the diagram above
     for the sake of simplicity.

   - The example depicted in Figure 1 shows the case where the Head-end
     and Tail-end areas are connected by means of area 0. The case of
     an inter-area TE LSP between two IGP areas that does not transit
     through area 0 is not precluded.

3.3. Example of Topology for the Inter-AS TE Case

   We consider the following general case, built on a superset of the
   various scenarios defined in [RFC4216]:

```
         <-- AS1 ----> <------- AS2 ------><--- AS3 ----->

                 <---BGP--->             <---BGP-->
     CE1---R0---X1-ASBR1-----ASBR4--R3---ASBR7----ASBR9----R6
           |\      \ |       / |   / |   / |           |       |
           | \      ASBR2---/ ASBR5  | -- |           |       |
           | \      |        |      |/    |           |       |
           R1-R2---ASBR3-----ASBR6--R4---ASBR8----ASBR10---R7---CE2

           <======= Inter-AS TE LSP (LSR to LSR)===========>
      or

      <======== Inter-AS TE LSP (CE to ASBR) =>

      or

      <================= Inter-AS TE LSP (CE to CE)===============>
```
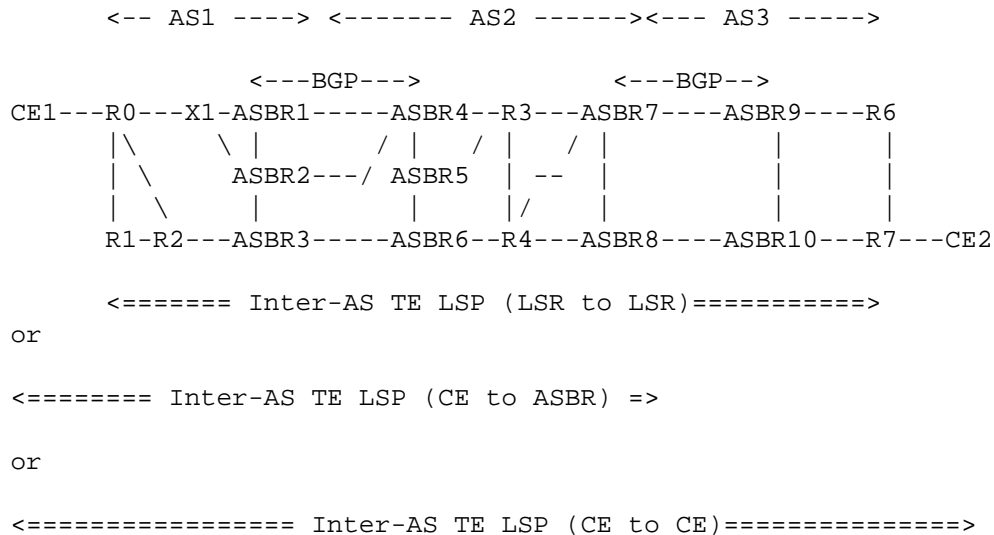
        Figure 2 - Example of topology for the inter-AS TE case

   The diagram depicted in Figure 2 covers all the inter-AS TE
   deployment cases described in [RFC4216].

   Description of Figure 2:

   - Three interconnected ASs, respectively AS1, AS2, and AS3. Note
     that in some scenarios described in [RFC4216] AS1=AS3.

   - The ASBRs in different ASs are BGP peers. There is usually no IGP
     running on the single hop links interconnecting the ASBRs and also
     referred to as inter-ASBR links.

   - Each AS runs an IGP (IS-IS or OSPF) with the required IGP TE
     extensions (see [RFC3630], [RFC3784], [RFC4203] and [RFC4205]). In
     other words, the ASs are TE enabled.

   - CE: Customer Edge router.

   - Each AS can be made of several IGP areas. The path computation
     technique described in this document applies to the case of a
     single AS made of multiple IGP areas, multiple ASs made of a single
     IGP area, or any combination of the above. For the sake of
     simplicity, each routing domain will be considered as a single area

     in this document. The case of an inter-AS TE LSP spanning multiple
     ASs where some of those ASs are themselves made of multiple IGP
     areas can be easily derived from the examples above: the per-domain
     path computation technique described in this document is applied
     recursively in this case.

     - An inter-AS TE LSP T1 originated at R0 in AS1 and terminated at R6
       in AS3.

4. Per-Domain Path Computation Procedures

   The mechanisms for inter-domain TE LSP computation as described in
   this document can be used regardless of the nature of the
   inter-domain TE LSP (contiguous, stitched, or nested).

   Note that any path can be defined as a set of loose and strict hops.
   In other words, in some cases, it might be desirable to rely on the
   dynamic path computation in some domains, and exert a strict control
   on the path in other domains (defining strict hops).

   When an LSR that is a boundary node such as an ABR/ASBR receives a
   Path message with an ERO that contains a strict node, the procedures
   specified in [RFC3209] apply and no further action is needed.

   When an LSR that is a boundary node such as an ABR/ASBR receives a
   Path message with an ERO that contains a loose hop or an abstract
   node that is not a simple abstract node (that is, an abstract node
   that identifies more than one LSR), then it MUST follow the
   procedures as described in [RFC5151].

   In addition, the following procedures describe the path computation
   procedures that SHOULD be carried out on the LSR:

   1) If the next hop is not present in the TED, the two following
      conditions MUST be checked:

      o Whether the IP address of the next-hop boundary LSR is outside
         of the current domain

      o Whether the next-hop domain is PSC (Packet Switch Capable) and
         uses in-band control channel

   If the two conditions above are satisfied, then the boundary LSR
   SHOULD check if the next hop has IP reachability (via IGP or BGP).
   If the next hop is not reachable, then a signaling failure occurs and
   the LSR SHOULD send back an RSVP PathErr message upstream with error
   code=24 ("Routing Problem") and error subcode as described in section
   4.3.4 of [RFC3209]. If the available routing information indicates

that next hop is reachable, the selected route will be expected to
pass through a domain boundary via a domain boundary LSR. The
determination of domain boundary point based on routing information
is what we term as "auto-discovery" in this document. In the absence
of such an auto-discovery mechanism, a) the ABR in the case of
inter-area TE or the ASBR in the next-hop AS in the case of inter-AS
TE should be the signaled loose next hop in the ERO and hence should
be accessible via the TED, or b) there needs to be an alternate
scheme that provides the domain exit points. Otherwise, the path
computation for the inter-domain TE LSP will fail.

An implementation MAY support the ability to disable such an IP
reachability fall-back option should the next-hop boundary LSR not be
present in the TED. In other words, an implementation MAY support
the possibility to trigger a signaling failure whenever the next hop
is not present in the TED.

2) Once the next-hop boundary LSR has been determined (according to
   the procedure described in 1)) or if the next-hop boundary is
   present in the TED:

   o Case of a contiguous TE LSP. Unless not allowed by policy, the
      boundary LSR that processes the ERO SHOULD perform an ERO
      expansion (a process consisting of computing the constrained
      path up to the next loose hop and adding the list of hops as
      strict nodes in the ERO). If no path satisfying the set of
      constraints can be found, then this is treated as a path
      computation and signaling failure and an RSVP PathErr message
      SHOULD be sent for the inter-domain TE LSP based on section
      4.3.4 of [RFC3209].

   o Case of a stitched or nested TE LSP

      * If the boundary LSR is a candidate LSR for intra-area H-LSP/
         S-LSP setup (the boundary has local policy for nesting or
         stitching), the TE LSP is a candidate for hierarchy/nesting
         (the "Contiguous LSP" bit defined in [RFC5151] is not set),
         and if there is no H-LSP/S-LSP from this LSR to the next-hop
         boundary LSR that satisfies the constraints, it SHOULD
         signal an H-LSP/S-LSP to the next-hop boundary LSR. If a
         pre-configured H-LSP(s) or S-LSP(s) already exists, then it
         will try to select from among those intra-domain LSPs.
         Depending on local policy, it MAY signal a new H-LSP/S-LSP
         if this selection fails. If the H-LSP/S-LSP is successfully
         signaled or selected, it propagates the inter-domain Path
         message to the next hop following the procedures described
         in [RFC5151]. If for some reason the dynamic H-LSP/S-LSP
         setup to the next-hop boundary LSR fails, then this SHOULD

be treated as a path computation and signaling failure and
an RSVP PathErr message SHOULD be sent upstream for the
inter-domain LSP. Similarly, if selection of a pre-
configured H-LSP/S-LSP fails and local policy prevents
dynamic H-LSP/S, this SHOULD be treated as a path
computation and signaling failure and an RSVP PathErr
message SHOULD be sent upstream for the inter-domain TE LSP.
In both of these cases, procedures described in section
4.3.4 of [RFC3209] SHOULD be followed to handle the failure.

*   If, however, the boundary LSR is not a candidate for
    intra-domain H-LSP/S-LSP (the boundary LSR does not have
    local policy for nesting or stitching) or the TE LSP is not
    a candidate for hierarchy/nesting (the "Contiguous LSP" bit
    defined in [RFC5151] is set), then it SHOULD apply the same
    procedure as for the contiguous case.

The ERO of an inter-domain TE LSP may comprise abstract nodes such as
ASs. In such a case, upon receiving the ERO whose next hop is an AS,
the boundary LSR has to determine the next-hop boundary LSR, which
may be determined based on the auto-discovery process mentioned
above. If multiple ASBR candidates exist, the boundary LSR may apply
some policies based on peering contracts that may have been
pre-negotiated. Once the next-hop boundary LSR has been determined,
a similar procedure as the one described above is followed.

Note the following related to the inter-AS TE case:

In terms of computation of an inter-AS TE LSP path, an interesting
optimization technique consists of allowing the ASBRs to flood the TE
information related to the inter-ASBR link(s) although no IGP TE is
enabled over those links (and so there is no IGP adjacency over the
inter-ASBR links). This of course implies that the inter-ASBR links
be TE-enabled although no IGP is running on those links.

```
        <-- AS1 ----> <------- AS2 ------><--- AS3 ----->


                <---BGP--->              <---BGP-->
      CE1---R0---X1-ASBR1-----ASBR4--R3---ASBR7----ASBR9----R6
           |\      \ |       / |    / |   / |            |        |
           | \       ASBR2---/ ASBR5  | -- |            |        |
           | \       |        |      |/    |            |        |
           R1-R2---ASBR3-----ASBR6--R4---ASBR8----ASBR10---R7---CE2
```

        Figure 3 - Flooding of the TE-related information for
                   the inter-ASBR links

Referring to Figure 3, ASBR1 for example would advertise in its OSPF
Link State Advertisement (LSA)/IS-IS LSP the Traffic Engineering TLVs
related to the link ASBR1-ASBR4.

This allows an LSR (could be the entry ASBR) in the previous AS to
make a more appropriate route selection up to the entry ASBR in the
immediately downstream AS taking into account the constraints
associated with the inter-ASBR links. This reduces the risk of call
setup failure due to inter-ASBR links not satisfying the inter-AS TE
LSP set of constraints. Note that the TE information is only related
to the inter-ASBR links: the TE LSA/LSP flooded by the ASBR includes
not only the TE-enabled links contained in the AS but also the
inter-ASBR links.

Note that no summarized TE information is leaked between ASs, which
is compliant with the requirements listed in [RFC4105] and [RFC4216].

For example, consider the diagram depicted in Figure 2: when ASBR1
floods its IGP TE LSA ((opaque LSA for OSPF)/LSP (TLV 22 for IS-IS))
in its routing domain, it reflects the reservation states and TE
properties of the following links: X1-ASBR1, ASBR1-ASBR2, and
ASBR1-ASBR4.

Thanks to such an optimization, the inter-ASBR TE link information
corresponding to the links originated by the ASBR is made available
in the TED of other LSRs in the same domain to which the ASBR
belongs. Consequently, the path computation for an inter-AS TE LSP
path can also take into account the inter-ASBR link(s). This will
improve the chance of successful signaling along the next AS in case
of resource shortage or unsatisfied constraints on inter-ASBR links,
and it potentially reduces one level of crankback. Note that no
topology information is flooded, and these links are not used in IGP
SPF computations. Only the TE information for the outgoing links
directly connected to the ASBR is advertised.

Note that an operator may decide to operate a stitched segment or
1-hop hierarchical LSP for the inter-ASBR link.

4.1. Example with an Inter-Area TE LSP

The following example uses Figure 1 as a reference.

4.1.1. Case 1: T0 Is a Contiguous TE LSP

The Head-end LSR (R0) first determines the next-hop ABR (which could
be manually configured by the user or dynamically determined by using
the auto-discovery mechanism). R0 then computes the path to reach
the selected next-hop ABR (ABR1) and signals the Path message. When

   the Path message reaches ABR1, it first determines the next-hop ABR
   from its area 0 along the LSP path (say, ABR3), either directly from
   the ERO (if for example the next-hop ABR is specified as a loose hop
   in the ERO) or by using the auto-discovery mechanism specified above.

   - Example 1 (set of loose hops):
     R0-ABR1(loose)-ABR3(loose)-R1(loose)

   - Example 2 (mix of strict and loose hops):
     R0-X1-ABR1-ABR3(loose)-X2-X3-R1

   Note that a set of paths can be configured on the Head-end LSR,
   ordered by priority. Each priority path can be associated with a
   different set of constraints. It may be desirable to systematically
   have a last-resort option with no constraint to ensure that the
   inter-area TE LSP could always be set up if at least a TE path exists
   between the inter-area TE LSP source and destination. In case of
   setup failure or when an RSVP PathErr is received indicating that the
   TE LSP has suffered a failure, an implementation might support the
   possibility of retrying a particular path option a configurable
   amount of times (optionally with dynamic intervals between each
   trial) before trying a lower-priority path option.

   Once it has computed the path up to the next-hop ABR (ABR3), ABR1
   sends the Path message along the computed path. Upon receiving the
   Path message, ABR3 then repeats a similar procedure. If ABR3 cannot
   find a path obeying the set of constraints for the inter-area TE LSP,
   the signaling process stops and ABR3 sends a PathErr message to ABR1.
   Then ABR1 can in turn trigger a new path computation by selecting
   another egress boundary LSR (ABR4 in the example above) if crankback
   is allowed for this inter-area TE LSP (see [RFC4920]). If crankback
   is not allowed for that inter-area TE LSP or if ABR1 has been
   configured not to perform crankback, then ABR1 MUST stop the
   signaling process and MUST forward a PathErr up to the Head-end LSR
   (R0) without trying to select another ABR.

4.1.2. Case 2: T0 Is a Stitched or Nested TE LSP

   The Head-end LSR (R0) first determines the next-hop ABR (which could
   be manually configured by the user or dynamically determined by using
   the auto-discovery mechanism). R0 then computes the path to reach
   the selected next-hop ABR and signals the Path message. When the
   Path message reaches ABR1, it first determines the next-hop ABR from
   its area 0 along the LSP path (say ABR3), either directly from the
   ERO (if for example the next-hop ABR is specified as a loose hop in
   the ERO) or by using an auto-discovery mechanism, specified above.

ABR1 then checks whether it has an H-LSP or S-LSP to ABR3 matching
the constraints carried in the inter-area TE LSP Path message. If
not, ABR1 computes the path for an H-LSP or S-LSP from ABR1 to ABR3
satisfying the constraint and sets it up accordingly. Note that the
H-LSP or S-LSP could have also been pre-configured.

Once ABR1 has selected the H-LSP/S-LSP for the inter-area LSP, using
the signaling procedures described in [RFC5151], ABR1 sends the Path
message for the inter-area TE LSP to ABR3. Note that irrespective of
whether ABR1 does nesting or stitching, the Path message for the
inter-area TE LSP is always forwarded to ABR3. ABR3 then repeats the
exact same procedures. If ABR3 cannot find a path obeying the set of
constraints for the inter-area TE LSP, ABR3 sends a PathErr message
to ABR1. Then ABR1 can in turn either select another H-LSP/S-LSP to
ABR3 if such an LSP exists or select another egress boundary LSR
(ABR4 in the example above) if crankback is allowed for this inter-
area TE LSP (see [RFC4920]). If crankback is not allowed for that
inter-area TE LSP or if ABR1 has been configured not to perform
crankback, then ABR1 forwards the PathErr up to the inter-area Head-
end LSR (R0) without trying to select another egress LSR.

4.2. Example with an Inter-AS TE LSP

The following example uses Figure 2 as a reference.

The path computation procedures for establishing an inter-AS TE LSP
are very similar to those of an inter-area TE LSP described above.
The main difference is related to the presence of inter-ASBR link(s).

4.2.1. Case 1: T1 Is a Contiguous TE LSP

The inter-AS TE path may be configured on the Head-end LSR as a set
of strict hops, loose hops, or a combination of both.

- Example 1 (set of loose hops):
  ASBR4(loose)-ASBR9(loose)-R6(loose)

- Example 2 (mix of strict and loose hops):
  R2-ASBR3-ASBR2-ASBR1-ASBR4-ASBR10(loose)-ASBR9-R6

In example 1 above, a per-AS path computation is performed,
respectively on R0 for AS1, ASBR4 for AS2, and ASBR9 for AS3. Note
that when an LSR has to perform an ERO expansion, the next hop either
must belong to the same AS or must be the ASBR directly connected to
the next hop AS. In this latter case, the ASBR reachability is
announced in the IGP TE LSA/LSP originated by its neighboring ASBR.
In example 1 above, the TE LSP path is defined as: ASBR4(loose)-
ASBR9(loose)-R6(loose). This implies that R0 must compute the path

   from R0 to ASBR4, hence the need for R0 to get the TE reservation
   state related to the ASBR1-ASBR4 link (flooded in AS1 by ASBR1). In
   addition, ASBR1 must also announce the IP address of ASBR4 specified
   in T1's path configuration.

   Once it has computed the path up to the next-hop ASBR, ASBR1 sends
   the Path message for the inter-area TE LSP to ASBR4 (supposing that
   ASBR4 is the selected next-hop ASBR). ASBR4 then repeats the exact
   same procedures. If ASBR4 cannot find a path obeying the set of
   constraints for the inter-AS TE LSP, then ASBR4 sends a PathErr
   message to ASBR1. Then ASBR1 can in turn either select another ASBR
   (ASBR5 in the example above) if crankback is allowed for this inter-
   AS TE LSP (see [RFC4920]), or if crankback is not allowed for that
   inter-AS TE LSP or if ASBR1 has been configured not to perform
   crankback, ABR1 stops the signaling process and forwards a PathErr up
   to the Head-end LSR (R0) without trying to select another egress LSR.
   In this case, the Head-end LSR can in turn select another sequence of
   loose hops, if configured. Alternatively, the Head-end LSR may
   decide to retry the same path; this can be useful in case of setup
   failure due to an outdated IGP TE database in some downstream AS. An
   alternative could also be for the Head-end LSR to retry the same
   sequence of loose hops after having relaxed some constraint(s).

4.2.2. Case 2: T1 Is a Stitched or Nested TE LSP

   The path computation procedures are very similar to the inter-area
   LSP setup case described earlier. In this case, the H-LSPs or S-LSPs
   are originated by the ASBRs at the entry to the AS.

5. Path Optimality/Diversity

   Since the inter-domain TE LSP is computed on a per-domain (area, AS)
   basis, one cannot guarantee that the optimal inter-domain path can be
   found.

   Moreover, computing two diverse paths using a per-domain path
   computation approach may not be possible in some topologies (due to
   the well-known "trapping" problem).

   For example, consider the following simple topology:

```
                      +-------+
                     /         \
                    A----B-----C------D
                         \         /
                          +--------+
```

              Figure 4 - Example of the "trapping" problem

In the simple topology depicted in Figure 4, with a serialized
approach using the per-domain path computation technique specified in
this document, a first TE LSP may be computed following the path
A-B-C-D, in which case no diverse path could be found although two
diverse paths actually exist: A-C-D and A-B-D. The aim of that
simple example that can easily be extended to the inter-domain case
is to illustrate the potential issue of not being able to find
diverse paths using the per-domain path computation approach when
diverse paths exist.

As already pointed out, the required path computation method can be
selected by the Service Provider on a per-LSP basis.

If the per-domain path computation technique does not meet the set of
requirements for a particular TE LSP (e.g., path optimality,
requirements for a set of diversely routed TE LSPs), other techniques
such as PCE-based path computation techniques may be used (see
[RFC4655]).

6. Reoptimization of an Inter-Domain TE LSP

As stated in [RFC4216] and [RFC4105], the ability to reoptimize an
already established inter-domain TE LSP constitutes a requirement.
The reoptimization process significantly differs based upon the
nature of the TE LSP and the mechanism in use for the TE LSP
computation.

The following mechanisms can be used for reoptimization and are
dependent on the nature of the inter-domain TE LSP.

6.1. Contiguous TE LSPs

After an inter-domain TE LSP has been set up, a better route might
appear within any traversed domain. Then in this case, it is
desirable to get the ability to reroute an inter-domain TE LSP in a
non-disruptive fashion (making use of the so-called Make-Before-Break
procedure) to follow a better path. This is a known as a TE LSP
reoptimization procedure.

[RFC4736] proposes a mechanism that allows the Head-end LSR to be
notified of the existence of a more optimal path in a downstream
domain. The Head-end LSR may then decide to gracefully reroute the
TE LSP using the Make-Before-Break procedure. In case of a
contiguous LSP, the reoptimization process is strictly controlled by
the Head-end LSR that triggers the Make-Before-Break procedure as
defined in [RFC3209], regardless of the location of the better path.

6.2. Stitched or Nested (non-contiguous) TE LSPs

   In the case of a stitched or nested inter-domain TE LSP, the
   reoptimization process is treated as a local matter to any domain.
   The main reason is that the inter-domain TE LSP is a different LSP
   (and therefore different RSVP session) from the intra-domain S-LSP or
   H-LSP in an area or an AS. Therefore, reoptimization in a domain is
   done by locally reoptimizing the intra-domain H-LSP or S-LSP. Since
   the inter-domain TE LSPs are transported using S-LSP or H-LSP across
   each domain, optimality of the inter-domain TE LSP in a domain is
   dependent on the optimality of the corresponding S-LSP or H-LSP. If
   after an inter-domain LSP is set up a more optimal path is available
   within a domain, the corresponding S-LSP or H-LSP will be reoptimized
   using Make-Before-Break techniques discussed in [RFC3209].
   Reoptimization of the H-LSP or S-LSP automatically reoptimizes the
   inter-domain TE LSPs that the H-LSP or S-LSP transports.
   Reoptimization parameters like frequency of reoptimization, criteria
   for reoptimization like metric or bandwidth availability, etc. can
   vary from one domain to another and can be configured as required,
   per intra-domain TE S-LSP or H-LSP if it is pre-configured or based
   on some global policy within the domain.

   Hence, in this scheme, since each domain takes care of reoptimizing
   its own S-LSPs or H-LSPs, and therefore the corresponding
   inter-domain TE LSPs, the Make-Before-Break can happen locally and is
   not triggered by the Head-end LSR for the inter-domain LSP. So, no
   additional RSVP signaling is required for LSP reoptimization, and
   reoptimization is transparent to the Head-end LSR of the inter-domain
   TE LSP.

   If, however, an operator desires to manually trigger reoptimization
   at the Head-end LSR for the inter-domain TE LSP, then this solution
   does not prevent that. A manual trigger for reoptimization at the
   Head-end LSR SHOULD force a reoptimization thereby signaling a "new"
   path for the same LSP (along the more optimal path) making use of the
   Make-Before-Break procedure. In response to this new setup request,
   the boundary LSR either may initiate new S-LSP setup, in case the
   inter-domain TE LSP is being stitched to the intra-domain S-LSP, or
   it may select an existing or new H-LSP, in case of nesting. When the
   LSP setup along the current path is complete, the Head-end LSR should
   switch over the traffic onto that path, and the old path is
   eventually torn down. Note that the Head-end LSR does not know a
   priori whether a more optimal path exists. Such a manual trigger
   from the Head-end LSR of the inter-domain TE LSP is, however, not
   considered to be a frequent occurrence.

      Procedures described in [RFC4736] MUST be used if the operator does
      not desire local reoptimization of certain inter-domain LSPs. In
      this case, any reoptimization event within the domain MUST be
      reported to the Head-end node. This SHOULD be a configurable policy.

6.3. Path Characteristics after Reoptimization

      Note that in the case of loose hop reoptimization of contiguous
      inter-domain TE LSP or local reoptimization of stitched/nested S-LSP
      where boundary LSRs are specified as loose hops, the TE LSP may
      follow a preferable path within one or more domain(s) but would still
      traverse the same set of boundary LSRs. In contrast, in the case of
      PCE-based path computation techniques, because the end-to-end optimal
      path is computed, the reoptimization process may lead to following a
      completely different inter-domain path (including a different set of
      boundary LSRs).

7. Security Considerations

      Signaling of inter-domain TE LSPs raises security issues (discussed
      in section 7 of [RFC5151]).

      [RFC4726] provides an overview of the requirements for security in an
      MPLS-TE or GMPLS multi-domain environment. In particular, when
      signaling an inter-domain RSVP-TE LSP, an operator may make use of
      the security features already defined for RSVP-TE ([RFC3209]). This
      may require some coordination between the domains to share the keys
      (see [RFC2747] and [RFC3097]), and care is required to ensure that
      the keys are changed sufficiently frequently. Note that this may
      involve additional synchronization, should the domain border nodes be
      protected with Fast Reroute ([RFC4090], since the Merge Point (MP)
      and Point of Local Repair (PLR) should also share the key. For an
      inter-domain TE LSP, especially when it traverses different
      administrative or trust domains, the following mechanisms SHOULD be
      provided to an operator (also see [RFC4216]):

      1) A way to enforce policies and filters at the domain borders to
         process the incoming inter-domain TE LSP setup requests (Path
         messages) based on certain agreed trust and service
         levels/contracts between domains. Various LSP attributes such as
         bandwidth, priority, etc. could be part of such a contract.

      2) A way for the operator to rate-limit LSP setup requests or error
         notifications from a particular domain.

      3) A mechanism to allow policy-based outbound RSVP message processing
         at the domain border node, which may involve filtering or
         modification of certain addresses in RSVP objects and messages.

This document relates to inter-domain path computation. It must be
noted that the process for establishing paths described in this
document does not increase the information exchanged between ASs and
preserves topology confidentiality, in compliance with [RFC4105] and
[RFC4216]. That being said, the signaling of inter-domain TE LSP
according to the procedure defined in this document requires path
computation on boundary nodes that may be exposed to various attacks.
Thus, it is RECOMMENDED to support policy decisions to reject the ERO
expansion of an inter-domain TE LSP if not allowed.

8. Acknowledgements

   We would like to acknowledge input and helpful comments from Adrian
   Farrel, Jean-Louis Le Roux, Dimitri Papadimitriou, and Faisal Aslam.

9. References

9.1. Normative References

   [RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate
                Requirement Levels", BCP 14, RFC 2119, March 1997.

   [RFC3209]   Awduche, D., Berger, L., Gan, D., Li, T., Srinivasan, V.,
                and G. Swallow, "RSVP-TE: Extensions to RSVP for LSP
                Tunnels", RFC 3209, December 2001.

   [RFC3473]   Berger, L., Ed., "Generalized Multi-Protocol Label
                Switching (GMPLS) Signaling Resource ReserVation
                Protocol-Traffic Engineering (RSVP-TE) Extensions", RFC
                3473, January 2003.

9.2. Informative References

   [RFC4920]   Farrel, A., Ed., Satyanarayana, A., Iwata, A., Fujita,
                N., and G. Ash, "Crankback Signaling Extensions for MPLS
                and GMPLS RSVP-TE", RFC 4920, July 2007.

   [RFC5151]   Farrel, A., Ed., Ayyangar, A., and JP. Vasseur, "Inter-
                Domain MPLS and GMPLS Traffic Engineering -- Resource
                Reservation Protocol-Traffic Engineering (RSVP-TE)
                Extensions", RFC 5151, February 2008.

   [RFC5150]   Ayyangar, A., Kompella, K., Vasseur, JP., and A. Farrel,
                "Label Switched Path Stitching with Generalized
                Multiprotocol Label Switching Traffic Engineering (GMPLS
                TE)", RFC 5150, February 2008.

   [RFC2702]   Awduche, D., Malcolm, J., Agogbua, J., O'Dell, M., and J.
               McManus, "Requirements for Traffic Engineering Over
               MPLS", RFC 2702, September 1999.

   [RFC2747]   Baker, F., Lindell, B., and M. Talwar, "RSVP
               Cryptographic Authentication", RFC 2747, January 2000.

   [RFC3097]   Braden, R. and L. Zhang, "RSVP Cryptographic
               Authentication -- Updated Message Type Value", RFC 3097,
               April 2001.

   [RFC3630]   Katz, D., Kompella, K., and D. Yeung, "Traffic
               Engineering (TE) Extensions to OSPF Version 2", RFC 3630,
               September 2003.

   [RFC3784]   Smit, H. and T. Li, "Intermediate System to Intermediate
               System (IS-IS) Extensions for Traffic Engineering (TE)",
               RFC 3784, June 2004.

   [RFC4090]   Pan, P., Ed., Swallow, G., Ed., and A. Atlas, Ed., "Fast
               Reroute Extensions to RSVP-TE for LSP Tunnels", RFC 4090,
               May 2005.

   [RFC4105]   Le Roux, J.-L., Ed., Vasseur, J.-P., Ed., and J. Boyle,
               Ed., "Requirements for Inter-Area MPLS Traffic
               Engineering", RFC 4105, June 2005.

   [RFC4203]   Kompella, K., Ed., and Y. Rekhter, Ed., "OSPF Extensions
               in Support of Generalized Multi-Protocol Label Switching
               (GMPLS)", RFC 4203, October 2005.

   [RFC4205]   Kompella, K., Ed., and Y. Rekhter, Ed., "Intermediate
               System to Intermediate System (IS-IS) Extensions in
               Support of Generalized Multi-Protocol Label Switching
               (GMPLS)", RFC 4205, October 2005.

   [RFC4216]   Zhang, R., Ed., and J.-P. Vasseur, Ed., "MPLS Inter-
               Autonomous System (AS) Traffic Engineering (TE)
               Requirements", RFC 4216, November 2005.

   [RFC4655]   Farrel, A., Vasseur, J.-P., and J. Ash, "A Path
               Computation Element (PCE)-Based Architecture", RFC 4655,
               August 2006.

   [RFC4726]   Farrel, A., Vasseur, J.-P., and A. Ayyangar, "A Framework
               for Inter-Domain Multiprotocol Label Switching Traffic
               Engineering", RFC 4726, November 2006.

   [RFC4736]   Vasseur, JP., Ed., Ikejiri, Y., and R. Zhang,
                "Reoptimization of Multiprotocol Label Switching (MPLS)
                Traffic Engineering (TE) Loosely Routed Label Switched
                Path (LSP)", RFC 4736, November 2006.

Authors' Addresses

   JP Vasseur (editor)
   Cisco Systems, Inc.
   1414 Massachusetts Avenue
   Boxborough, MA 01719
   USA

   EMail: jpv@cisco.com

   Arthi Ayyangar (editor)
   Juniper Networks
   1194 N. Mathilda Ave
   Sunnyvale, CA 94089
   USA

   EMail: arthi@juniper.net

   Raymond Zhang
   BT
   2160 E. Grand Ave.
   El Segundo, CA 90025
   USA

   EMail: raymond.zhang@bt.com

Full Copyright Statement

Intellectual Property