                        Distribution of Diverse BGP Paths

Abstract

   The BGP4 protocol specifies the selection and propagation of a single
   best path for each prefix. As defined and widely deployed today, BGP
   has no mechanisms to distribute alternate paths that are not
   considered best path between its speakers. This behavior results in
   a number of disadvantages for new applications and services.

   The main objective of this document is to observe that by simply
   adding a new session between a route reflector and its client, the
   Nth best path can be distributed. This document also compares
   existing solutions and proposed ideas that enable distribution of
   more paths than just the best path.

   This proposal does not specify any changes to the BGP protocol
   definition. It does not require a software upgrade of provider edge
   (PE) routers acting as route reflector clients.

Status of This Memo

Table of Contents

1. Introduction

   The current BGP4 protocol specification [RFC4271] allows for the
   selection and propagation of only one best path for each prefix. As
   defined today, the BGP protocol has no mechanism to distribute paths
   other than best path between its speakers. This behavior results in
   a number of problems in the deployment of new applications and
   services.

This document presents a mechanism for solving the problem based on
the conceptual creation of parallel route-reflector planes. It also
compares existing solutions and proposes ideas that enable
distribution of more paths than just the best path. The parallel
route-reflector planes solution brings very significant benefits at a
negligible capex and opex deployment price as compared to the
alternative techniques (full BGP mesh or add-paths [ADD-PATHS]) and
is being considered by a number of network operators for deployment
in their networks.

This proposal does not specify any changes to the BGP protocol
definition. It does not require upgrades to provider edge or core
routers, nor does it need network-wide upgrades. The only upgrade
required is the new functionality on the new or current route
reflectors.

2. History

The need to disseminate more paths than just the best path is
primarily driven by three issues. The first is the problem of BGP
oscillations [RFC3345]. The second is the desire for faster
reachability restoration in the event of failure of the network link
or network element. The third is a need to enhance BGP load-
balancing capabilities. These issues have led to the proposal of BGP
add-paths [ADD-PATHS].

2.1. BGP Add-Paths Proposal

As it has been proven that distribution of only the best path of a
route is not sufficient to meet the needs of the continuously growing
number of services carried over BGP, the add-paths proposal was
submitted in 2002 to enable BGP to distribute more than one path.
This is achieved by including an additional four-octet value called
the "Path Identifier" as a part of the Network Layer Reachability
Information (NLRI).

The implication of this change on a BGP implementation is that it
must now maintain a per-path, instead of per-prefix, peer
advertisement state to track to which of the peers a given path was
advertised. This new requirement comes with its own memory and
processing cost.

An important observation is that distribution of more than one best
path by the Autonomous System Border Routers (ASBRs) with multiple
External BGP (EBGP) peers attached where no "next-hop self" is set
may result in inconsistent best-path selection within the autonomous
system. Therefore, it is also required to attach the possible
tiebreakers in the form of a new attribute and propagate those within

the domain. The example of such an attribute for the purpose of fast
connectivity restoration to address that very case of ASBR injecting
multiple external paths into the Internal BGP (IBGP) mesh has been
presented and discussed in "Advertisement of Multiple Paths in BGP"
[ADD-PATHS]. Based on the additionally propagated information, best-
path selection is recommended to be modified to make sure that best-
and backup-path selection within the domain stays consistent. More
discussion on this particular point is contained in Section 6,
"Deployment Considerations". In the proposed solution in this
document, we observe that to address most of the applications, just
use of the best external advertisement is required. For ASBRs that
are peering to multiple upstream domains, setting "next-hop self" is
recommended.

The add-paths protocol extensions have to be implemented by all the
routers within an Autonomous System (AS) in order for the system to
work correctly. Analyzing the benefits or risks associated with
partial add-paths deployments remains quite a topic for research.
The risk becomes even greater in networks not using some form of
edge-to-edge encapsulation.

The required code modifications can offer the foundation for
enhancements, such as the "Fast Connectivity Restoration Using BGP
Add-path" [FAST-CONN]. The deployment of such technology in an
entire service-provider network requires software, and perhaps
sometimes, in the case of End-of-Engineering or End-of-Life
equipment, even hardware upgrades. Such an operation may or may not
be economically feasible. Even if add-path functionality was
available today on all commercial routing equipment and across all
vendors, experience indicates that it may easily take years to
achieve 100% deployment coverage within any medium or large global
network.

While it needs to be clearly acknowledged that the add-path mechanism
provides the most general way to address the problem of distributing
many paths between BGP speakers, this document provides a solution
that is much easier to deploy and requires no modification to the BGP
protocol where only a few additional paths may be required. The
alternative method presented is capable of addressing critical
service-provider requirements for disseminating more than a single
path across an AS with a significantly lower deployment cost. That,
in light of the number of general network scaling concerns documented
in RFC 4984 [RFC4984], "Report from the IAB Workshop on Routing and
Addressing", may provide a significant advantage.

3. Goals

   The proposal described in this document is not intended to compete
   with add-paths. It provides an interim solution until add-paths are
   standardized and implemented and until support for that function can
   be deployed across the network.

   It is presented to network operators as a possible choice and
   provides those operators who need additional paths today an
   alternative from the need to transition to a full mesh. The Nth best
   path describes a set of N paths with different BGP next hops with no
   implication of ordering or preference among said N paths.

   It is intended as a way to buy more time, allowing for a smoother and
   gradual migration where router upgrades will be required for,
   perhaps, different reasons. It will also allow the time required so
   that standard RP/RE memory size can easily accommodate the associated
   overhead with other techniques without any compromises.

4. Multi-Plane Route Reflection

   The idea contained in the proposal assumes the use of route
   reflection within the network.

   Let's observe today's picture of a simple route-reflected domain:

```
                              ASBR3
                               ***
                              *   *
           +------------*   *-----------+
           | AS1        *   *           |
           |             ***            |
           |                            |
           |                            |
           |                            |
           | RR1         ***        RR2 |
           | ***        *   *       *** |
           |*   *       * P *      *   *|
           |*   *       *   *      *   *|
           | ***         ***       ***  |
           |                            |
           |            IBGP            |
           |                            |
           |                            |
           |     ***          ***       |
           |    *   *        *   *       |
           +-----*   *---------*   *----+
                *   *          *   *
                 ***            ***
              ASBR1          ASBR2
                       EBGP
              Figure 1: Simple route reflection
```

   Abbreviations used:
      RR - Route Reflector
      P - Core router

   Figure 1 shows an AS that is connected via EBGP peering at ASBR1 and
   ASBR2 to an upstream AS or set of ASes. For a given destination "D",
   ASBR1 and ASBR2 may have an external path P1 and P2, respectively.
   The AS network uses two route reflectors, RR1 and RR2, for redundancy
   reasons. The route reflectors propagate the single BGP best path for
   each route to all clients. All ASBRs are clients of RR1 and RR2.

Following are the possible cases of the path information that ASBR3
may receive from route reflectors RR1 and RR2:

1. When the best-path tiebreaker is the IGP distance: When paths P1
    and P2 are considered to be equally good best-path candidates,
    the selection will depend on the distance of the path's next hops
    from the route reflector making the decision. Depending on the
    positioning of the route reflectors in the IGP topology, they may
    choose the same best path or a different one. In such a case,
    ASBR3 may receive either the same path or different paths from
    each of the route reflectors.

2. When the best-path tiebreaker is MULTI_EXIT_DISC (MED) or
    LOCAL_PREF: In this case, only one path from the preferred exit
    point ASBR will be available to RRs since the other peering ASBR
    will consider the IBGP path as best and will not announce (or if
    already announced will withdraw) its own external path. The
    exception here is the use of the BGP Best-External proposal
    [EXT-PATH], which will allow a stated ASBR to still propagate to
    the RRs on its own external path. Unfortunately, RRs will not be
    able to distribute it any further to other clients, as only the
    overall best path will be reflected.

There is no requirement of path ordering. The "Nth best path" really
describes set of N paths with different BGP next hops.

The proposed solution is based on the use of additional route
reflectors or new functionality enabled on the existing route
reflectors that, instead of distributing the best path for each
route, will distribute an alternative path other than best. The
best-path (main) reflector plane distributes the best path for each
route as it does today. The second plane distributes the second best
path for each route, and so on. Distribution of N paths for each
route can be achieved by using N reflector planes.

As diverse-path functionality may be enabled on a per-peer basis, one
of the deployment models can be realized to continue advertisement of
the overall best path from both route reflectors, while in addition a
new session can be provisioned to get an additional path. This will
allow the uninterrupted use of the best path, even if one of the RRs
goes down, provided that the overall best path is still a valid one.

Each plane of the route reflectors is a logical entity and may or may
not be co-located with the existing best-path route reflectors.
Adding a route-reflector plane to a network may be as easy as
enabling a logical router partition, new BGP process, or just a new
configuration knob on an existing route reflector and configuring an
additional IBGP session from the current clients if required. There

are no code changes required on the route-reflector clients for this
mechanism to work. It is easy to observe that the installation of
one or more additional route-reflector control planes is much cheaper
and is easier than upgrading hundreds of route-reflector clients in
the entire network to support different BGP protocol encoding.

Diverse-path route reflectors need the new ability to calculate and
propagate the Nth best path instead of the overall best path. An
implementation is encouraged to enable this new functionality on a
per-neighbor basis.

While this is an implementation detail, the code to calculate the Nth
best path is also required by other BGP solutions. For example, in
the application of fast connectivity restoration, BGP must calculate
a backup path for installation into the Routing Information Base
(RIB) and Forwarding Information Base (FIB) ahead of the actual
failure.

To address the problem of external paths not being available to route
reflectors due to LOCAL_PREF or MED factors, it is recommended that
ASBRs enable [EXT-PATH] functionality in order to always inject their
external paths to the route reflectors.

4.1. Co-located Best- and Backup-Path RRs

To simplify the description, let's assume that we only use two route-
reflector planes (N=2). When co-located, the additional second-best-
path reflectors are connected to the network at the same points from
the perspective of the IGP as the existing best-path RRs. Let's also
assume that best-external functionality is enabled on all ASBRs.

```
                            ASBR3
                             ***
                            *   *
          +------------*   *-----------+
          | AS1          *   *         |
          |               ***          |
          |                            |
          | RR1                   RR2  |
          |  ***                  ***  |
          | *    *      ***      *    *|
          | *    *     *   *     *    *|
          |  ***      * P *      ***   |
          | *    *     *   *     *    *|
          | *    *      ***      *    *|
          |  ***                  ***  |
          | RR1'       IBGP       RR2' |
          |                            |
          |                            |
          |     ***          ***       |
          |    *   *        *   *      |
          +-----*   *---------*   *----+
                *   *        *   *
                 ***          ***
              ASBR1          ASBR2


                            EBGP
```

         Figure 2: Co-located Second-Best-Path RR Plane

   The following is a list of configuration changes required to enable
   the second-best-path route-reflector plane:

   1. Unless the same RR1/RR2 platform is being used, adding RR1' and
       RR2' either as the logical or physical new control-plane RRs in
       the same IGP points as RR1 and RR2, respectively.

   2. Enabling best-external functionality on ASBRs.

   3. Enabling RR1' and RR2' for second plane route reflection.
       Alternatively, instructing existing RR1 and RR2 to calculate the
       second-best path also.

   4. Unless one of the existing RRs is set to advertise only diverse
       path to its current clients, configuring new ASBRs-RR' IBGP
       sessions.

The expected behavior is that under any BGP condition, the ASBR3 and
P routers will receive both paths P1 and P2 for destination D. The
availability of both paths will allow them to implement a number of
new services as listed in Section 8 ("Applications").

As an alternative to fully meshing all RRs and RRs', an operator that
has a large number of reflectors deployed today may choose to peer
newly introduced RRs' to a hierarchical RR', which would be an IBGP
interconnect point within the second plane as well as between planes.

One deployment model of this scenario can be achieved by simply
upgrading the existing route reflectors without deploying any new
logical or physical platforms. Such an upgrade would allow route
reflectors to service both peers that have upgraded to add-paths, as
well as those peers that cannot be immediately upgraded while at the
same time allowing distribution of more than a single best path. The
obvious protocol benefit of using existing RRs to distribute towards
their clients' best and diverse BGP paths over different IBGP
sessions is the automatic assurance that such a client would always
get different paths with their next hop being different.

The way to accomplish this would be to create a separate IBGP session
for each Nth BGP path. Such a session should be preferably
terminated at a different loopback address of the route reflector.
At the BGP OPEN stage of each such session, a different bgp_router_id
may be used. Correspondingly, the route reflector should also allow
its clients to use the same bgp_router_id on each such session.

4.2. Randomly Located Best- and Backup-Path RRs

Now let's consider a deployment case in which an operator wishes to
enable a second RR' plane using only a single additional router in a
different network location from his current route reflectors. This
model would be of particular use in networks in which some form of
end-to-end encapsulation (IP or MPLS) is enabled between provider-
edge routers.

Note that this model of operation assumes that the present best-path
route reflectors are only control-plane devices. If the route
reflector is in the data-forwarding path, then the implementation
must be able to clearly separate the Nth best-path selection from the
selection of the paths to be used for data forwarding. The basic
premise of this mode of deployment assumes that all reflector planes
have the same information to choose from, which includes the same set
of BGP paths. It also requires the ability to ignore the step of
comparison of the IGP metric to reach the BGP next hop during best-
path calculation.

```
                           ASBR3
                            ***
                           *   *
          +-----------*   *-----------+
          | AS1          *   *        |
          | IBGP         ***          |
          |                           |
          |             ***           |
          |            *   *          |
          | RR1        * P *     RR2   |
          | ***        *   *     ***  |
          |*   *        ***      *   *|
          |*   *                 *   *|
          | ***        RR'        ***  |
          |            ***            |
          |           *   *           |
          |           *   *           |
          |            ***            |
          |    ***           ***      |
          |   *   *         *   *     |
          +-----*   *---------*   *----+
              *   *         *   *
               ***           ***
            ASBR1          ASBR2


                          EBGP

   Figure 3: Experimental Deployment of Second-Best-Path RR Plane
```

   The following is a list of configuration changes required to enable
   the second-best-path route reflector RR' as a single platform or to
   enable one of the existing control-plane RRs for diverse-path
   functionality:

   1. If needed, adding RR' logical or physical as a new route
       reflector anywhere in the network.

   2. Enabling best-external functionality on ASBRs.

   3. Disabling IGP metric check in BGP best path on all route
       reflectors.

   4. Enabling RR' or any of the existing RR for second plane path
       calculation.

   5. If required, fully meshing newly added RRs' with all the other
       reflectors in both planes. This condition does not apply if the
       newly added RR'(s) already have peering to all ASBRs/PEs.

   6. Configure new BGP sessions between ASBRs and RRs (unless one of
       the existing RRs is set to advertise only diverse path to its
       current clients).

   In this scenario, the operator has the flexibility to introduce the
   new additional route-reflector functionality on any existing or new
   hardware in the network. Any existing routers that are not already
   members of the best-path route-reflector plane can be easily
   configured to serve the second plane either by using a
   logical/virtual router partition or by having their BGP
   implementation compliant to this specification.

   Even if the IGP metric is not taken into consideration when comparing
   paths during the best-path calculation, an implementation still has
   to consider paths with unreachable next hops invalid. It is worth
   pointing out that some implementations today already allow for
   configuration that results in no IGP metric comparison during the
   best-path calculation.

   The additional planes of route reflectors do not need to be fully
   redundant as the primary plane does. If we are preparing for a
   single network failure event, a failure of a non-backed-up Nth best-
   path route reflector would not result in a connectivity outage of the
   actual data plane. The reason is that this would, at most, affect
   the presence of a backup path (not an active one) on the same parts
   of the network. If the operator chooses to create the Nth best-path
   plane redundantly by installing not one, but two or more route
   reflectors serving each additional plane, the additional robustness
   will be achieved.

   As a result of this solution, ASBR3 and other ASBRs peering to RR'
   will be receiving the second best path.

   Similarly to Section 4.1, as an alternative to fully meshing all RRs
   and diverse path RRs', operators may choose to peer newly introduced
   RRs' to a hierarchical RR', which would be an IBGP interconnect point
   between planes.

   It is recommended that an implementation advertise the overall best
   path over the Nth diverse-path session if there is no other BGP path
   with a different next hop present. This is equivalent to today's
   case where the client is connected to more than one RR.

4.3. Multi-Plane Route Servers for Internet Exchanges

   Another group of devices in which the proposed multi-plane
   architecture may be of particular applicability is the EBGP route
   servers used at many Internet exchange points.

In such cases, hundreds of ISPs are interconnected on a common LAN.
Instead of having hundreds of direct EBGP sessions on each exchange
client, a single peering is created to the transparent route server.
The route server can only propagate a single best path. Mandating
the upgrade for hundreds of different service providers in order to
implement add-path may be much more difficult as compared to asking
them to provision one new EBGP session to an Nth best path route
server plane. This allows the distribution of more than the single
best BGP path from a given route server to such an Internet exchange
point (IX) peer.

The solution proposed in this document fits very well with the
requirement of having broader EBGP path diversity among the members
of any Internet exchange point.

5. Discussion on Current Models of IBGP Route Distribution

   In today's networks, BGP4 operates as specified in [RFC4271].

   There are a number of technology choices for intra-AS BGP route
   distribution:

   1. Full mesh

   2. Confederations

   3. Route reflectors

5.1. Full Mesh

   A full mesh, the most basic IBGP architecture, exists when all BGP
   speaking routers within the AS peer directly with all other BGP
   speaking routers within the AS, irrespective of where a given router
   resides within the AS (e.g., P router, PE router, etc.).

   While this is the simplest intra-domain path-distribution method,
   historically, there have been a number of challenges in realizing
   such an IBGP full mesh in a large-scale network. While some of these
   challenges are no longer applicable, the following (as well as
   others) may still apply:

   1. Number of TCP sessions: The number of IBGP sessions on a single
       router in a full-mesh topology of a large-scale service provider
       can easily reach hundreds. Such numbers could be a concern on
       hardware and software used in the late 70s, 80s, and 90s. Today,
       customer requirements for the number of BGP sessions per box are
       reaching thousands. This is already an order of magnitude more
       than the potential number of IBGP sessions. Advancements in the

hardware and software used in production routers means that
running a full mesh of IBGP sessions should not be dismissed due
to the resulting number of TCP sessions alone.

2. Provisioning: When operating and troubleshooting large networks,
   one of the topmost requirements is to keep the design as simple
   as possible. When the autonomous system's network is composed of
   hundreds of nodes, it becomes very difficult to manually
   provision a full mesh of IBGP sessions. Adding or removing a
   router requires reconfiguration of all other routers in the AS.
   While this is a real concern today, there is already work in
   progress in the IETF to define IBGP peering automation through an
   IBGP Auto Discovery mechanism [AUTO-MESH].

3. Number of paths: Another concern when deploying a full IBGP mesh
   is the number of BGP paths for each route that have to be stored
   at every node. This number is very tightly related to the number
   of external peerings of an AS, the use of LOCAL_PREF or MED
   techniques, and the presence of best-external [EXT-PATH]
   advertisement configuration. If we make a rough assumption that
   the BGP4-path data structure consumes about 80-100 bytes, the
   resulting control-plane memory requirement for 500,000 IPv4
   routes with one additional external path is 38-48 MB, while for 1
   million IPv4 routes, it grows linearly to 76-95 MB. It is not
   possible to reach a general conclusion if this condition is
   negligible or if it is a show stopper for a full-mesh deployment
   without direct reference to a given network.

To summarize, a full-mesh IBGP peering can offer natural
dissemination of multiple external paths among BGP speakers. When
realized with the help of IBGP Auto Discovery peering automation,
this seems like a viable deployment, especially in medium- and small-
scale networks.

5.2. Confederations

For the purpose of this document, let's observe that confederations
[RFC5065] can be viewed as a hierarchical full-mesh model.

Within each sub-AS, BGP speakers are fully meshed, and as discussed
in Section 2.1, all full-mesh characteristics (number of TCP
sessions, provisioning, and potential concern over number of paths
still apply in the sub-AS scale).

In addition to the direct peering of all BGP speakers within each
sub-AS, all sub-AS border routers must also be fully meshed with each
other. Sub-AS border routers configured with best-external
functionality can inject additional (diverse) paths within a sub-AS.

To summarize, it is technically sound to use confederations with the
combination of best-external to achieve distribution of more than a
single best path per route in a large autonomous systems.

In topologies where route reflectors are deployed within the
confederation sub-ASes, the technique described here applies.

5.3. Route Reflectors

The main motivation behind the use of route reflectors [RFC4456] is
the avoidance of the full-mesh session management problem described
above. Route reflectors, for good or for bad, are the most common
solution today for interconnecting BGP speakers within an internal
routing domain.

Route-reflector peerings follow the advertisement rules defined by
the BGP4 protocol. As a result, only a single best path per prefix
is sent to client BGP peers. This is the main reason many current
networks are exposed to a phenomenon called BGP path starvation,
which essentially results in the inability to deliver a number of
applications discussed later.

When interconnecting BGP speakers between domains, the route
reflection equivalent is popularly called the "Route Server" and is
globally deployed today in many Internet exchange points.

6. Deployment Considerations

Distribution of the diverse-BGP-paths proposal allows the
dissemination of more paths than just the best path to the route-
reflector or route-server clients of today's BGP4 implementations.
As a deployment recommendation, it needs to be mentioned that fast
connectivity restoration as well as a majority of intra-domain BGP-
level load balancing needs can be accommodated with only two paths
(overall best and second best). Therefore, as a deployment
recommendation, this document suggests use of N=2 with diverse-path.

From the client's point of view, receiving additional paths via
separate IBGP sessions terminated at the new route-reflector plane is
functionally equivalent to constructing a full-mesh peering without
the problems such a full mesh would come with, as discussed in
earlier section.

By precisely defining the number of reflector planes, network
operators have full control over the number of redundant paths in the
network. This number can be defined to address the needs of the
service(s) being deployed.

The Nth-plane route reflectors should act as control-plane network
entities. While they can be provisioned on the current production
routers, selected Nth-best BGP paths should not be used directly in
the date plane with the exception of such paths being BGP multipath
eligible and such functionality is enabled. Regarding RRs being in
the data plane unless multipath is enabled, the second best path is
expected to be a backup path and should be installed as such into the
local RIB/FIB.

The use of the term "planes" in this document is more of a conceptual
nature. In practice, all paths are still kept in the single table
where normal best path is calculated. This means that tools like the
looking glass should not observe any changes or impact when
diverse-path has been enabled.

The proposed architecture deployed along with the BGP best-external
functionality covers all three cases where the classic BGP route-
reflection paradigm would fail to distribute alternate (diverse)
paths. These are

1. ASBRs advertising their single best-external paths with no
    LOCAL_PREF or MED present.

2. ASBRs advertising their single best-external paths with
    LOCAL_PREF or MED present and with BGP best-external
    functionality enabled.

3. ASBRs with multiple external paths.

This section focuses on discussion of case 3 above in more detail.
This describes the scenario of a single ASBR connected to multiple
EBGP peers. In practice, this peering scenario is quite common. It
is mostly due to the geographic location of EBGP peers and the
diversity of those peers (for example, peering to multiple tier-1
ISPs, etc.). It is not designed for failure-recovery scenarios, as
single failure of the ASBR would simultaneously result in loss of
connectivity to all of the peers. In most medium and large
geographically distributed networks, there is always another ASBR or
multiple ASBRs providing peering backups, typically in other
geographically diverse locations in the network.

When an operator uses ASBRs with multiple peerings, setting next-hop
self will effectively allow local repair of the atomic failure of any
external peer without any compromise to the data plane.
Traditionally, the most common reason for not setting next-hop self
is the associated drawback of losing the ability to signal the
external failures of peering ASBRs or links to those ASBRs by fast
IGP flooding. Such a potential drawback can be easily avoided by
using a different peering address from the address used for next-hop
mapping and removing the next-hop from the IGP at the last possible
BGP path failure.

Herein, one may correctly observe that in the case of setting next-
hop self on an ASBR, attributes of other external paths such that the
ASBR is peering with may be different from the attributes of its best
external path. Therefore, not injecting all of those external paths
with their corresponding attributes cannot be compared to equivalent
paths for the same prefix coming from different ASBRs.

While such observation, in principle, is correct, one should put
things in perspective of the overall goal, which is to provide data-
plane connectivity upon a single failure with minimal
interruption/packet loss. During such transient conditions, using
even potentially suboptimal exit points is reasonable, so long as
forwarding information loops are not introduced. In the mean time,
the BGP control plane will on its own re-advertise the newly elected
best external path, and route-reflector planes will calculate their
Nth best paths and propagate them to its clients. The result is that
after seconds, even if potential suboptimality were encountered, it
will be quickly and naturally healed.

7. Summary of Benefits

   Distribution of the diverse-BGP-paths proposal provides the following
   benefits when compared to the alternatives:

   1. No modifications to the BGP4 protocol.

   2. No requirement for upgrades to edge and core routers (as required
       in [ADD-PATHS]). It is backward compatible with the existing BGP
       deployments.

   3. Can be easily enabled by the introduction of a new route
       reflector, a route server plane dedicated to the selection and
       distribution of Nth best-path, or just by new configuration of
       the upgraded current route reflector(s).

   4. Does not require major modification to BGP implementations in the
      entire network, which would result in an unnecessary increase of
      memory and CPU consumption due to the shift from today's per-
      prefix to a per-path advertisement state tracking.

   5. Can be safely deployed gradually on an RR cluster basis.

   6. The proposed solution is equally applicable to any BGP address
      family as described in "Multiprotocol Extensions for BGP-4"
      [RFC4760]. In particular, it can be used "as is" without any
      modifications to both IPv4 and IPv6 address families.

8. Applications

   This section lists the most common applications that require the
   presence of redundant BGP paths:

   1. Fast connectivity restoration in which backup paths with
      alternate exit points would be pre-installed as well as
      pre-resolved in the FIB of routers. This allows for a local
      action upon reception of a critical event notification of
      network/node failure. This failure recovery mechanism that is
      based on the presence of backup paths is also suitable for
      gracefully addressing scheduled maintenance requirements as
      described in [BGP-SHUTDOWN].

   2. Multi-path load balancing for both IBGP and EBGP.

   3. BGP control-plane churn reduction for both intra-domain and
      inter-domain.

   An important point to observe is that all of the above intra-domain
   applications are based on the use of reflector planes but are also
   applicable in the inter-domain Internet exchange point examples. As
   discussed in Section 4.3, an Internet exchange can conceptually
   deploy shadow route server planes, each responsible for distribution
   of an Nth best path to its EBGP peers. In practice, it may just be
   equal to a new short configuration and establishment of new BGP
   sessions to IX peers.

9. Security Considerations

   The new mechanism for diverse BGP path dissemination proposed in this
   document does not introduce any new security concerns as compared to
   the base BGP4 specification [RFC4271] and especially when compared
   against full-IBGP-mesh topology.

   In addition, the authors observe that all BGP security issues as
   described in [RFC4272] apply to the additional BGP session or
   sessions as recommended by this specification. Therefore, all
   recommended mitigation techniques to BGP security are applicable
   here.

10. Contributors

   The following people contributed significantly to the content of the
   document:

   Selma Yilmaz
   Cisco Systems
   170 West Tasman Drive
   San Jose, CA 95134
   US
   Email: seyilmaz@cisco.com

   Satish Mynam
   Juniper Networks
   1194 N. Mathilda Ave
   Sunnyvale, CA 94089
   US
   Email: smynam@juniper.net

   Isidor Kouvelas
   Cisco Systems
   170 West Tasman Drive
   San Jose, CA 95134
   US
   Email: kouvelas@cisco.com

11. Acknowledgments

   The authors would like to thank Bruno Decraene, Bart Peirens, Eric
   Rosen, Jim Uttaro, Renwei Li, Wes George, and Adrian Farrel for their
   valuable input.

   The authors would also like to express a special thank you to a
   number of operators who helped optimize the provided solution to be
   as close as possible to their daily operational practices. In
   particular, many thanks to Ted Seely, Shane Amante, Benson
   Schliesser, and Seiichi Kawamura.

12. References

12.1. Normative References

   [RFC4271]   Rekhter, Y., Ed., Li, T., Ed., and S. Hares, Ed., "A
               Border Gateway Protocol 4 (BGP-4)", RFC 4271, January
               2006.

   [RFC4456]   Bates, T., Chen, E., and R. Chandra, "BGP Route
               Reflection: An Alternative to Full Mesh Internal BGP
               (IBGP)", RFC 4456, April 2006.

   [RFC4760]   Bates, T., Chandra, R., Katz, D., and Y. Rekhter,
               "Multiprotocol Extensions for BGP-4", RFC 4760, January
               2007.

   [RFC5226]   Narten, T. and H. Alvestrand, "Guidelines for Writing an
               IANA Considerations Section in RFCs", BCP 26, RFC 5226,
               May 2008.

12.2. Informative References

   [ADD-PATHS] Walton, D., Chen, E., Retana, A., and J. Scudder,
               "Advertisement of Multiple Paths in BGP", Work in
               Progress, June 2012.

   [AUTO-MESH] Raszuk, R., "IBGP Auto Mesh", Work in Progress, January
               2004.

   [BGP-SHUTDOWN]
               Decraene, B., Francois, P., Pelsser, C., Ahmad, Z., and
               A. Armengol, "Requirements for the Graceful Shutdown of
               BGP Sessions", Work in Progress, September 2009.

   [EXT-PATH] Marques, P., Fernando, R., Chen, E., Mohapatra, P., and
               H. Gredler, "Advertisement of the Best External Route in
               BGP", Work in Progress, January 2012.

   [FAST-CONN] Mohapatra, P., Fernando, R., Filsfils, C., and R. Raszuk,
               "Fast Connectivity Restoration Using BGP Add-path", Work
               in Progress), October 2011.

   [RFC3345]   McPherson, D., Gill, V., Walton, D., and A. Retana,
               "Border Gateway Protocol (BGP) Persistent Route
               Oscillation Condition", RFC 3345, August 2002.

   [RFC4272]   Murphy, S., "BGP Security Vulnerabilities Analysis", RFC
               4272, January 2006.

   [RFC5065]   Traina, P., McPherson, D., and J. Scudder, "Autonomous
               System Confederations for BGP", RFC 5065, August 2007.

Authors' Addresses

   Robert Raszuk (editor)
   NTT MCL
   101 S Ellsworth Avenue Suite 350
   San Mateo, CA 94401
   United States

   EMail: robert@raszuk.net

   Rex Fernando
   Cisco Systems
   170 West Tasman Drive
   San Jose, CA 95134
   United States

   EMail: rex@cisco.com

   Keyur Patel
   Cisco Systems
   170 West Tasman Drive
   San Jose, CA 95134
   United States

   EMail: keyupate@cisco.com

   Danny McPherson
   Verisign, Inc.
   12061 Bluemont Way
   Reston, VA 20190
   United States

   EMail: dmcpherson@verisign.com

   Kenji Kumaki
   KDDI Corporation
   Garden Air Tower
   Iidabashi, Chiyoda-ku, Tokyo 102-8460
   Japan

   EMail: ke-kumaki@kddi.com