                Remote Procedure Call (RPC) Security Version 3

Abstract

   This document specifies version 3 of the Remote Procedure Call (RPC)
   security protocol (RPCSEC_GSS).  This protocol provides support for
   multi-principal authentication of client hosts and user principals to
   a server (constructed by generic composition), security label
   assertions for multi-level security and type enforcement, structured
   privilege assertions, and channel bindings.  This document updates
   RFC 5403.

Status of This Memo

   This is an Internet Standards Track document.

   This document is a product of the Internet Engineering Task Force
   (IETF).  It represents the consensus of the IETF community.  It has
   received public review and has been approved for publication by the
   Internet Engineering Steering Group (IESG).  Further information on
   Internet Standards is available in Section 2 of RFC 7841.

   Information about the current status of this document, any errata,
   and how to provide feedback on it may be obtained at
   http://www.rfc-editor.org/info/rfc7861.

Table of Contents

1.  Introduction and Motivation

   The original Remote Procedure Call (RPC) security protocol
   (RPCSEC_GSS) [RFC2203] provided for authentication of RPC clients and
   servers to each other using the Generic Security Service Application
   Programming Interface (GSS-API) [RFC2743].  The second version of
   RPCSEC_GSS [RFC5403] added support for channel bindings [RFC5056].

   Existing GSS-API mechanisms are insufficient for communicating
   certain authorization and authentication information to a server.
   The GSS-API and its mechanisms certainly could be extended to address
   this shortcoming.  However, it is addressed here at the application
   layer, i.e., in RPCSEC_GSS.

   A major motivation for version 3 of RPCSEC_GSS (RPCSEC_GSSv3) is to
   add support for multi-level (labeled) security and server-side copy
   for NFSv4.

Multi-Level Security (MLS) is a traditional model where subjects
(processes) are given a security level (Unclassified, Secret,
Top Secret, etc.) and objects (files) are given security labels that
mandate the access of the subject to the object (see Section 9.1 of
[RFC7862]).

Labeled NFS (see Section 9 of [RFC7862]) uses an MLS policy with
Mandatory Access Control (MAC) systems as defined in [RFC4949].
Labeled NFS stores MAC file object labels on the NFS server and
enables client Guest Mode MAC as described in Section 9.5.3 of
[RFC7862].  RPCSEC_GSSv3 label assertions assert client MAC process
subject labels to enable Full Mode MAC when combined with Labeled NFS
as described in Section 9.5.1 of [RFC7862].

A traditional inter-server file copy entails the user gaining access
to a file on the source, reading it, and writing it to a file on the
destination.  In secure NFSv4 inter-server server-side copy (see
Section 4 of [RFC7862]), the user first secures access to both source
and destination files and then uses NFSv4.2-defined RPCSEC_GSSv3
structured privileges to authorize the destination to copy the file
from the source on behalf of the user.

Multi-principal authentication can be used to address shared cache
poisoning attacks (see Section 9 of [AFS-RXGK]) on the client cache
by a user.  As described in Section 7 of [AFS-RXGK], multi-user
machines with a single cache manager can fetch and cache data on a
user's behalf and re-display it for another user from the cache
without refetching the data from the server.  The initial data
acquisition is authenticated by the first user's credentials, and if
only that user's credentials are used, it may be possible for a
malicious user or users to "poison" the cache for other users by
introducing bogus data into the cache.

Another use of the multi-principal assertion is the secure conveyance
of privilege information for processes running with more (or even
with less) privilege than the user normally would be accorded.

1.1.  Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
"SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
document are to be interpreted as described in RFC 2119 [RFC2119].

1.2.  Added Functionality

   RPCSEC_GSS version 3 (RPCSEC_GSSv3) is the same as RPCSEC_GSSv2
   [RFC5403], except that the following assertions of authority have
   been added:

   o  Security labels for Full Mode security type enforcement, and other
      labeled security models (see Section 9.5.1 of [RFC7862]).

   o  Application-specific structured privileges.  These allow an RPC
      application client to pass structured information to the
      corresponding application code in a server to control the use of
      the privilege and/or the conditions in which the privilege may be
      exercised.  For an example, see server-side copy as described in
      [RFC7862].

   o  Multi-principal authentication of the client host and user to the
      server, done by binding two RPCSEC_GSS handles.

   o  Simplified channel binding.

   Assertions of labels and privileges are evaluated by the server,
   which may then map the asserted values to other values, all according
   to server-side policy.  See [RFC7862].

   An option for enumerating server-supported Label Format Specifiers
   (LFSs) is provided.  See Section 9.1 of [RFC7862].

   Note that there is no RPCSEC_GSS_CREATE payload that is REQUIRED to
   implement.  RPCSEC_GSSv3 implementations are feature driven.  Besides
   implementing the RPCSEC_GSS_CREATE operation and payloads for the
   desired features, all RPCSEC_GSSv3 implementations MUST implement:

   o  The new RPCSEC_GSS version number (Section 2.2).

   o  The new reply verifier (Section 2.3).

   o  The new auth_stat values (Section 2.6).

RPCSEC_GSSv3 targets implementing a desired feature MUST also
implement the RPCSEC_GSS_LIST operation, and the RPCSEC_GSS_CREATE
operation replies for unsupported features as follows:

o  For label assertions, the target indicates no support by returning
   the new RPCSEC_GSS_LABEL_PROBLEM auth_stat value (see
   Section 2.7.1.3).

o  For structured privilege assertions, the target indicates no
   support by returning the new RPCSEC_GSS_UNKNOWN_MESSAGE auth_stat
   value (see Section 2.7.1.4).

o  For multi-principal authentication (Section 2.7.1.1), the target
   indicates no support by not including an rgss3_gss_mp_auth value
   in the rgss3_create_res.

o  For channel bindings (Section 2.7.1.2), the target indicates no
   support by not including an rgss3_chan_binding value in the
   rgss3_create_res.

1.3.  XDR Code Extraction

   This document contains the External Data Representation (XDR)
   [RFC4506] definitions for the RPCSEC_GSSv3 protocol.  The XDR
   description is provided in this document in a way that makes it
   simple for the reader to extract it into a form that is ready to
   compile.  The reader can feed this document in the following shell
   script to produce the machine-readable XDR description of
   RPCSEC_GSSv3:

   <CODE BEGINS>

   #!/bin/sh
   grep "^  *///" | sed 's?^  */// ??' | sed 's?^  *///$??'

   <CODE ENDS>

   That is, if the above script is stored in a file called "extract.sh"
   and this document is in a file called "spec.txt", then the reader
   can do:

   <CODE BEGINS>

    sh extract.sh < spec.txt > rpcsec_gss_v3.x

   <CODE ENDS>

The effect of the script is to remove leading white space from each
line, plus a sentinel sequence of "///".

2.  The RPCSEC_GSSv3 Protocol

RPCSEC_GSS version 3 (RPCSEC_GSSv3) is very similar to RPCSEC_GSS
version 2 (RPCSEC_GSSv2) [RFC5403].  The difference is that the new
support for assertions and channel bindings is implemented via a
different mechanism.

The entire RPCSEC_GSSv3 protocol is not presented here.  Only the
differences between RPCSEC_GSSv3 and RPCSEC_GSSv2 are shown.

RPCSEC_GSSv3 is implemented as follows:

o  A client uses an existing RPCSEC_GSSv3 context handle established
   in the usual manner (see Section 5.2 of [RFC2203]) to protect
   RPCSEC_GSSv3 exchanges; this will be termed the "parent" handle.

o  The server issues a "child" RPCSEC_GSSv3 handle in the
   RPCSEC_GSS_CREATE response, which uses the underlying GSS-API
   security context of the parent handle in all subsequent exchanges
   that use the child handle.

o  An RPCSEC_GSSv3 child handle MUST NOT be used as the parent handle
   in an RPCSEC_GSS3_CREATE control message.

2.1.  Compatibility with RPCSEC_GSSv2

The functionality of RPCSEC_GSSv2 [RFC5403] is fully supported by
RPCSEC_GSSv3, with the exception of the RPCSEC_GSS_BIND_CHANNEL
operation, which is not supported when RPCSEC_GSSv3 is in use (see
Section 2.5).

2.2.  Version Negotiation

An initiator that supports version 3 of RPCSEC_GSS simply issues an
RPCSEC_GSS request with the rgc_version field set to
RPCSEC_GSS_VERS_3.  If the target does not recognize
RPCSEC_GSS_VERS_3, the target will return an RPC error per
Section 5.1 of [RFC2203].

The initiator MUST NOT attempt to use an RPCSEC_GSS handle returned
by version 3 of a target with version 1 or version 2 of the same
target.  The initiator MUST NOT attempt to use an RPCSEC_GSS handle
returned by version 1 or version 2 of a target with version 3 of the
same target.

2.3.  New Reply Verifier

   A new reply verifier is needed for RPCSEC_GSSv3 because of a
   situation that arises from the use of the same GSS context by child
   and parent handles.  Because the RPCSEC_GSSv3 child handle uses the
   same GSS context as the parent handle, a child and parent
   RPCSEC_GSSv3 handle could have the same RPCSEC_GSS sequence numbers.
   Since the reply verifier of previous versions of RPCSEC_GSS computes
   a Message Integrity Code (MIC) on just the sequence number, this
   provides opportunities for man-in-the-middle attacks.

   This issue is addressed in RPCSEC_GSS version 3 by computing the
   verifier using exactly the same input as the information used to
   compute the request verifier, except that the mtype is changed from
   CALL to REPLY.  The new reply verifier computes a MIC over the
   following RPC reply header data:

     unsigned int xid;
     msg_type mtype;     /* set to REPLY */
     unsigned int rpcvers;
     unsigned int prog;
     unsigned int vers;
     unsigned int proc;
     opaque_auth  cred; /* binds the RPCSEC_GSS handle */

2.4.  XDR Code Preliminaries

   The following code fragment replaces the corresponding preliminary
   code shown in Figure 1 of [RFC5403].  The values in the code fragment
   in Section 2.6 are additions to the auth_stat enumeration.
   Subsequent code fragments are additions to the code for version 2
   that support the new procedures defined in version 3.

   <CODE BEGINS>

       ///   /*
       ///    * Copyright (c) 2016 IETF Trust and the persons
       ///    * identified as the authors.  All rights reserved.
       ///    *
       ///    * The authors of the code are identified in RFC 2203,
       ///    * RFC 5403, and RFC 7861.
       ///    *
       ///    * Redistribution and use in source and binary forms,
       ///    * with or without modification, are permitted
       ///    * provided that the following conditions are met:
       ///    *
       ///    * o Redistributions of source code must retain the above
       ///    *   copyright notice, this list of conditions and the
       ///    *   following disclaimer.
       ///    *
       ///    * o Redistributions in binary form must reproduce the
       ///    *   above copyright notice, this list of
       ///    *   conditions and the following disclaimer in
       ///    *   the documentation and/or other materials
       ///    *   provided with the distribution.
       ///    *
       ///    * o Neither the name of Internet Society, IETF or IETF
       ///    *   Trust, nor the names of specific contributors, may be
       ///    *   used to endorse or promote products derived from this
       ///    *   software without specific prior written permission.
       ///    *
       ///    *   THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS
       ///    *   AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED
       ///    *   WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE
       ///    *   IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS
       ///    *   FOR A PARTICULAR PURPOSE ARE DISCLAIMED.  IN NO
       ///    *   EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE
       ///    *   LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL,
       ///    *   EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT
       ///    *   NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR
       ///    *   SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS
       ///    *   INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF
       ///    *   LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY,

```
///   *   OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING
///   *    IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF
///   *    ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
///   */
///
///   /*
///    * This code was derived from RFC 2203, RFC 5403,
///    * and RFC 7861.  Please reproduce this note if possible.
///    */
///
///   enum rpc_gss_service_t {
///           /* Note: The enumerated value for 0 is reserved. */
///           rpc_gss_svc_none       = 1,
///           rpc_gss_svc_integrity  = 2,
///           rpc_gss_svc_privacy    = 3,
///           rpc_gss_svc_channel_prot = 4
///   };
///
///   enum rpc_gss_proc_t {
///           RPCSEC_GSS_DATA          = 0,
///           RPCSEC_GSS_INIT          = 1,
///           RPCSEC_GSS_CONTINUE_INIT = 2,
///           RPCSEC_GSS_DESTROY       = 3,
///           RPCSEC_GSS_BIND_CHANNEL  = 4, /* Not used */
///           RPCSEC_GSS_CREATE        = 5, /* New */
///           RPCSEC_GSS_LIST          = 6  /* New */
///   };
///
///   struct rpc_gss_cred_vers_1_t {
///           rpc_gss_proc_t   gss_proc; /* Control procedure */
///           unsigned int     seq_num;  /* Sequence number */
///           rpc_gss_service_t service;  /* Service used */
///           opaque           handle<>; /* Context handle */
///   };
///
///   const RPCSEC_GSS_VERS_1 = 1;
///   const RPCSEC_GSS_VERS_2 = 2;
///   const RPCSEC_GSS_VERS_3 = 3; /* New */
///
///   union rpc_gss_cred_t switch (unsigned int rgc_version) {
///   case RPCSEC_GSS_VERS_1:
///   case RPCSEC_GSS_VERS_2:
///   case RPCSEC_GSS_VERS_3: /* New */
///           rpc_gss_cred_vers_1_t rgc_cred_v1;
///   };
///

   <CODE ENDS>
```

   As seen above, the RPCSEC_GSSv3 credential has the same format as the
   RPCSEC_GSSv1 [RFC2203] and RPCSEC_GSSv2 [RFC5403] credential.
   Setting the rgc_version field to 3 indicates that the initiator and
   target support the new RPCSEC_GSSv3 control procedures.

2.5.  RPCSEC_GSS_BIND_CHANNEL Operation

   RPCSEC_GSSv3 provides a channel-binding assertion that replaces the
   RPCSEC_GSSv2 RPCSEC_GSS_BIND_CHANNEL operation.

   The RPCSEC_GSS_BIND_CHANNEL operation is not supported on RPCSEC_GSS
   version 3 handles.  If a server receives an RPCSEC_GSS_BIND_CHANNEL
   operation on an RPCSEC_GSSv3 handle, it MUST return a reply status of
   MSG_ACCEPTED with an accept_stat of PROC_UNAVAIL [RFC5531].

2.6.  New auth_stat Values

   RPCSEC_GSSv3 requires the addition of several values to the auth_stat
   enumerated type definition.  The use of these new auth_stat values is
   explained throughout this document.

```
            enum auth_stat {
                    ...
                    /*
                     * RPCSEC_GSSv3 errors
                     */
                    RPCSEC_GSS_INNER_CREDPROBLEM = 15,
                    RPCSEC_GSS_LABEL_PROBLEM     = 16,
                    RPCSEC_GSS_PRIVILEGE_PROBLEM = 17,
                    RPCSEC_GSS_UNKNOWN_MESSAGE   = 18
            };
```

2.7.  New Control Procedures

   There are two new RPCSEC_GSSv3 control procedures: RPCSEC_GSS_CREATE
   and RPCSEC_GSS_LIST.

   The RPCSEC_GSS_CREATE procedure binds any combination of assertions
   -- multi-principal authentication, labels, structured privileges, or
   channel bindings -- to a new RPCSEC_GSSv3 context returned in the
   rgss3_create_res rcr_handle field.

   The RPCSEC_GSS_LIST procedure queries the target for supported
   assertions.

RPCSEC_GSS version 3 control messages are similar to the RPCSEC_GSS
version 1 and version 2 RPCSEC_GSS_DESTROY control message (see
Section 5.4 of [RFC2203]) in that the sequence number in the request
must be valid and the header checksum in the verifier must be valid.
As in RPCSEC_GSS version 1 and version 2, the RPCSEC_GSS version 3
control messages may contain call data following the verifier in the
body of the NULLPROC procedure.  In other words, they look a lot like
an RPCSEC_GSS data message with the header procedure set to NULLPROC.

The client MUST use one of the following security services to protect
the RPCSEC_GSS_CREATE or RPCSEC_GSS_LIST control message:

o  rpc_gss_svc_integrity

o  rpc_gss_svc_privacy

Specifically, the client MUST NOT use rpc_gss_svc_none.

RPCSEC_GSS_LIST can also use rpc_gss_svc_channel_prot (see
RPCSEC_GSSv2 [RFC5403]) if the request is sent using an RPCSEC_GSSv3
child handle with channel bindings enabled as described in
Section 2.7.1.2.

2.7.1.  New Control Procedure - RPCSEC_GSS_CREATE

   <CODE BEGINS>

```
      ///  struct rgss3_create_args {
      ///          rgss3_gss_mp_auth    *rca_mp_auth;
      ///          rgss3_chan_binding   *rca_chan_bind_mic;
      ///          rgss3_assertion_u    rca_assertions<>;
      ///  };
      ///
      ///  struct rgss3_create_res {
      ///          opaque               rcr_handle<>;
      ///          rgss3_gss_mp_auth    *rcr_mp_auth;
      ///          rgss3_chan_binding   *rcr_chan_bind_mic;
      ///          rgss3_assertion_u    rcr_assertions<>;
      ///  };
      ///
      ///  enum rgss3_assertion_type {
      ///          LABEL = 0,
      ///          PRIVS = 1
      ///  };
      ///
      ///  union rgss3_assertion_u
      ///        switch (rgss3_assertion_type atype) {
      ///  case LABEL:
      ///          rgss3_label  rau_label;
      ///  case PRIVS:
      ///          rgss3_privs  rau_privs;
      ///  default:
      ///          opaque       rau_ext<>;
      ///  };
      ///
```

   <CODE ENDS>

   The call data for an RPCSEC_GSS_CREATE request consists of an
   rgss3_create_args, which binds one or more items of several kinds to
   the returned rcr_handle RPCSEC_GSSv3 context handle (the child
   handle):

   o  Multi-principal authentication: another RPCSEC_GSS context handle

   o  A channel binding

   o  Authorization assertions: labels and/or privileges

The reply to this message consists of either an error or an
rgss3_create_res structure.  As noted in Sections 2.7.1.3 and
2.7.1.4, successful rgss3_assertions are enumerated in rcr_assertions
and are REQUIRED to be enumerated in the same order as they appeared
in the rca_assertions argument.

Upon a successful RPCSEC_GSS_CREATE, both the client and the server
need to associate the resultant child rcr_handle context handle with
the parent context handle in their GSS context caches so as to be
able to reference the parent context given the child context handle.

RPCSEC_GSSv3 child handles MUST be destroyed upon the destruction of
the associated parent handle.

Server implementation and policy MAY result in labels, privileges,
and identities being mapped to concepts and values that are local to
the server.  Server policies should take into account the identity of
the client and/or user as authenticated via the GSS-API.

2.7.1.1.  Multi-Principal Authentication

   <CODE BEGINS>

      ///
      ///   struct rgss3_gss_mp_auth {
      ///           opaque          rgmp_handle<>; /* Inner handle */
      ///           opaque          rgmp_rpcheader_mic<>;
      ///   };
      ///

   <CODE ENDS>

RPCSEC_GSSv3 clients MAY assert a multi-principal authentication of
the RPC client host principal and a user principal.  This feature is
needed, for example, when an RPC client host wishes to use authority
assertions that the server may only grant if a user and an RPC client
host are authenticated together to the server.  Thus, a server may
refuse to grant requested authority to a user acting alone (e.g., via
an unprivileged user-space program) or to an RPC client host acting
alone (e.g., when an RPC client host is acting on behalf of a user)
but may grant requested authority to an RPC client host acting on
behalf of a user if the server identifies the user and trusts the RPC
client host.

It is assumed that an unprivileged user-space program would not have
access to RPC client host credentials needed to establish a GSS-API
security context authenticating the RPC client host to the server;
therefore, an unprivileged user-space program could not create an
RPCSEC_GSSv3 RPCSEC_GSS_CREATE message that successfully binds an RPC
client host and a user security context.

In addition to the parent handle (Section 2), the multi-principal
authentication call data has an RPCSEC_GSS version 3 handle
referenced via the rgmp_handle field termed the "inner" handle.
Clients using RPCSEC_GSSv3 multi-principal authentication MUST use an
RPCSEC_GSSv3 context handle that corresponds to a GSS-API security
context that authenticates the RPC client host for the parent handle.
The inner context handle of the multi-principal authentication
assertion MUST use an RPCSEC_GSSv3 context handle that corresponds to
a GSS-API security context that authenticates the user.  The reverse
(parent handle authenticates user, inner context handle authenticates
an RPC client host) MUST NOT be used.  Other multi-principal parent
and inner context handle uses might eventually make sense, but they
would need to be introduced in a new revision of the RPCSEC_GSS
protocol.

The child context handle returned by a successful multi-principal
assertion binds the inner RPCSEC_GSSv3 context handle to the parent
RPCSEC_GSS context handle and MUST be treated by servers as
authenticating the GSS-API initiator principal authenticated by the
inner context handle's GSS-API security context.  This principal may
be mapped to a server-side notion of user or principal.

Multi-principal binding is done by including an assertion of type
rgss3_gss_mp_auth in the RPCSEC_GSS_CREATE rgss3_create_args call
data.  The inner context handle is placed in the rgmp_handle field.
A MIC of the RPC header, up to and including the credential, is
computed using the GSS-API security context associated with the inner
context handle and is placed in the rgmp_rpcheader_mic field.  Note
that the rgmp_rpcheader_mic only identifies the client host GSS
context by its context handle (the parent context handle) in the RPC
header.

An RPCSEC_GSS_CREATE control procedure with a multi-principal
authentication payload MUST use the rpc_gss_svc_privacy security
service for protection.  This prevents an attacker from intercepting
the RPCSEC_GSS_CREATE control procedure, reassigning the (parent)
context handle, and stealing the user's identity.

The target verifies the multi-principal authentication by first
confirming that the parent context used is an RPC client host
context; the target then verifies the rgmp_rpcheader_mic using the
GSS-API security context associated with the rgmp_handle field.

On successful verification, the rgss3_gss_mp_auth field in the
rgss3_create_res reply MUST be filled in with the inner RPCSEC_GSSv3
context handle as the rgmp_handle and a MIC computed over the RPC
reply header (see Section 2.3) using the GSS-API security context
associated with the inner handle.

On failure, the rgss3_gss_mp_auth field is not sent
(rgss3_gss_mp_auth is an optional field).  A MSG_DENIED reply to the
RPCSEC_GSS_CREATE call is formulated as usual.

As described in Section 5.3.3.3 of [RFC2203], the server maintains a
list of contexts for the clients that are currently in session with
it.  When a client request comes in, there may not be a context
corresponding to its handle.  When this occurs on an
RPCSEC_GSS3_CREATE request processing of the parent handle, the
server rejects the request with a reply status of MSG_DENIED with the
reject_stat of AUTH_ERROR and with an auth_stat value of
RPCSEC_GSS_CREDPROBLEM.

A new value, RPCSEC_GSS_INNER_CREDPROBLEM, has been added to the
auth_stat type.  With a multi-principal authorization request, the
server must also have a context corresponding to the inner context
handle.  When the server does not have a context handle corresponding
to the inner context handle of a multi-principal authorization
request, the server sends a reply status of MSG_DENIED with the
reject_stat of AUTH_ERROR and with an auth_stat value of
RPCSEC_GSS_INNER_CREDPROBLEM.

When processing the multi-principal authentication request, if the
GSS_VerifyMIC() call on the rgmp_rpcheader_mic fails to return
GSS_S_COMPLETE, the server sends a reply status of MSG_DENIED with
the reject_stat of AUTH_ERROR and with an auth_stat value of
RPCSEC_GSS_INNER_CREDPROBLEM.

2.7.1.2.  Channel Binding

   <CODE BEGINS>

      ///
      ///  typedef opaque rgss3_chan_binding<>;
      ///

   <CODE ENDS>

   RPCSEC_GSSv3 provides a different way to do channel binding than
   RPCSEC_GSSv2 [RFC5403].  Specifically:

   a.  RPCSEC_GSSv3 builds on RPCSEC_GSSv1 by reusing existing,
       established context handles rather than providing a different RPC
       security flavor for establishing context handles.

   b.  Channel-bindings data is not hashed because there is now general
       agreement that it is the secure channel's responsibility to
       produce channel-bindings data of manageable size.

   (a) is useful in keeping RPCSEC_GSSv3 simple in general, not just for
   channel binding.  (b) is useful in keeping RPCSEC_GSSv3 simple
   specifically for channel binding.

   Channel binding is accomplished as follows.  The client prefixes the
   channel-bindings data octet string with the channel type as described
   in [RFC5056]; then, the client calls GSS_GetMIC() to get a MIC of the
   resulting octet string, using the parent RPCSEC_GSSv3 context
   handle's GSS-API security context.  The MIC is then placed in the
   rca_chan_bind_mic field of RPCSEC_GSS_CREATE arguments
   (rgss3_create_args).

   If the rca_chan_bind_mic field of the arguments of an
   RPCSEC_GSS_CREATE control message is set, then the server MUST verify
   the client's channel-binding MIC if the server supports this feature.
   If channel-binding verification succeeds, then the server MUST
   generate a new MIC of the same channel bindings and place it in the
   rcr_chan_bind_mic field of the RPCSEC_GSS_CREATE rgss3_create_res
   results.  If channel-binding verification fails or the server doesn't
   support channel binding, then the server MUST indicate this in its
   reply by not including an rgss3_chan_binding value in
   rgss3_create_res (rgss3_chan_binding is an optional field).

   The client MUST verify the result's rcr_chan_bind_mic value by
   calling GSS_VerifyMIC() with the given MIC and the channel-bindings
   data (including the channel-type prefix).  If client-side channel-
   binding verification fails, then the client MUST call

RPCSEC_GSS_DESTROY.  If the client requested channel binding but the
server did not include an rcr_chan_binding_mic field in the results,
then the client MAY continue to use the resulting context handle as
though channel binding had never been requested.  If the client
considers channel binding critical, it MUST call RPCSEC_GSS_DESTROY.

As per RPCSEC_GSSv2 [RFC5403]:

   Once a successful [channel-binding] procedure has been performed
   on an [RPCSEC_GSSv3] context handle, the initiator's
   implementation may map application requests for rpc_gss_svc_none
   and rpc_gss_svc_integrity to rpc_gss_svc_channel_prot credentials.
   And if the secure channel has privacy enabled, requests for
   rpc_gss_svc_privacy can also be mapped to
   rpc_gss_svc_channel_prot.

Any RPCSEC_GSSv3 child context handle that has been bound to a secure
channel in this way SHOULD be used only with the
rpc_gss_svc_channel_prot and SHOULD NOT be used with rpc_gss_svc_none
or rpc_gss_svc_integrity -- if the secure channel does not provide
privacy protection, then the client MAY use rpc_gss_svc_privacy where
privacy protection is needed or desired.

2.7.1.3.  Label Assertions

<CODE BEGINS>

```
///  struct rgss3_label {
///          rgss3_lfs       rl_lfs;
///          opaque          rl_label<>;
///  };
///
///  struct rgss3_lfs {
///          unsigned int rlf_lfs_id;
///          unsigned int rlf_pi_id;
///  };
///
```

<CODE ENDS>

The client discovers, via the RPCSEC_GSS_LIST control message, which
LFSs the server supports.  Full Mode MAC is enabled when an
RPCSEC_GSS version 3 process subject label assertion is combined with
a file object label provided by the NFSv4.2 sec_label attribute.

Label encoding is specified to mirror the NFSv4.2 sec_label attribute
described in Section 12.2.4 of [RFC7862].  The LFS is an identifier
used by the client to establish the syntactic format of the security

label and the semantic meaning of its components.  The Policy
Identifier (PI) is an optional part of the definition of an LFS that
allows clients and the server to identify specific security policies.
The opaque label field (rgss3_label) is dependent on the MAC model to
interpret and enforce.

If a label itself requires privacy protection (i.e., requires that
the user can assert that the label is a secret), then the client MUST
use the rpc_gss_svc_privacy protection service for the
RPCSEC_GSS_CREATE request.

RPCSEC_GSSv3 clients MAY assert a set of subject security labels in
some LFS by binding a label assertion to the RPCSEC_GSSv3 child
context handle.  This is done by including an assertion of type
rgss3_label in the RPCSEC_GSS_CREATE rgss3_create_args rca_assertions
call data.  The label assertion payload is the set of subject labels
asserted by the calling NFS client process.  The resultant child
context is used for NFS requests asserting the client process subject
labels.  The NFS server process that handles such requests then
asserts the (client) process subject label(s) as it attempts to
access a file that has associated Labeled NFS object labels.

Servers that support labeling in the requested LFS MAY map the
requested subject label to a different subject label as a result of
server-side policy evaluation.

The labels that are accepted by the target and bound to the
RPCSEC_GSSv3 context MUST be enumerated in the rcr_assertions field
of the rgss3_create_res RPCSEC_GSS_CREATE reply.

Servers that do not support labeling or that do not support the
requested LFS reject the label assertion with a reply status of
MSG_DENIED, a reject_status of AUTH_ERROR, and an auth_stat of
RPCSEC_GSS_LABEL_PROBLEM.

2.7.1.4.  Structured Privilege Assertions

   <CODE BEGINS>

```
///
///   typedef opaque utf8string<>;    /* UTF-8 encoding */
///   typedef utf8string utf8str_cs; /* Case-sensitive UTF-8 */
///
///   struct rgss3_privs {
///           utf8str_cs       rp_name<>;
///           opaque           rp_privilege<>;
///   };
```

   <CODE ENDS>

   A structured privilege is a capability defined by a specific RPC
   application.  To support the assertion of this privilege, by a client
   using the application, in a server that also supports the
   application, the application may define a private data structure that
   is understood by clients and servers implementing the RPC
   application.

   RPCSEC_GSSv3 clients MAY assert a structured privilege by binding the
   privilege to the RPCSEC_GSSv3 context handle.  This is done by
   including an assertion of type rgss3_privs in the RPCSEC_GSS_CREATE
   rgss3_create_args rca_assertions call data.

   The privilege is identified by the description string that is used by
   RPCSEC_GSSv3 to identify the privilege and communicate the private
   data between the relevant RPC application-specific code without
   needing to be aware of the details of the structure used.  Thus, as
   far as RPCSEC_GSSv3 is concerned, the defined structure is passed
   between client and server as opaque data encoded in the
   rpc_gss3_privs rp_privilege field.

   Encoding, server verification, and any server policies for structured
   privileges are described by the RPC application definition.  The
   rp_name field of rpc_gss3_privs carries the description string used
   to identify and list the privilege.  The utf8str_cs definition is
   from [RFC7530].

   A successful structured privilege assertion MUST be enumerated in the
   rcr_assertions field of the rgss3_create_res RPCSEC_GSS_CREATE reply.

   If a server receives a structured privilege assertion that it does
   not recognize, the assertion is rejected with a reply status of
   MSG_DENIED, a reject_status of AUTH_ERROR, and an auth_stat of
   RPCSEC_GSS_UNKNOWN_MESSAGE.

It is assumed that a client asserting more than one structured
privilege to be bound to a context handle would not require all the
privilege assertions to succeed.

The server MUST NOT reject RPCSEC_GSS_CREATE requests containing
supported structured privilege assertions, even if some of those
assertions are rejected (e.g., for local policy reasons).

If a server receives an RPCSEC_GSS_CREATE request containing one or
more unsupported structured privilege assertions, the request MUST be
rejected with a reply status of MSG_DENIED, a reject_status of
AUTH_ERROR, and an auth_stat of RPCSEC_GSS_PRIVILEGE_PROBLEM.

Section 4.9.1.1 of [RFC7862] ("Inter-Server Copy via ONC RPC with
RPCSEC_GSSv3") shows an example of structured privilege definition
and use.

2.7.2.  New Control Procedure - RPCSEC_GSS_LIST

   <CODE BEGINS>

```
    ///  enum rgss3_list_item {
    ///          LABEL = 0,
    ///          PRIVS = 1
    ///  };
    ///
    ///  struct rgss3_list_args {
    ///          rgss3_list_item     rla_list_what<>;
    ///  };
    ///
    ///  union rgss3_list_item_u
    ///        switch (rgss3_list_item itype) {
    ///  case LABEL:
    ///        rgss3_label         rli_labels<>;
    ///  case PRIVS:
    ///        rgss3_privs         rli_privs<>;
    ///  default:
    ///        opaque              rli_ext<>;
    ///  };
    ///
    ///  typedef rgss3_list_item_u rgss3_list_res<>;
    ///
```

   <CODE ENDS>

The call data for an RPCSEC_GSS_LIST request consists of a list of
integers (rla_list_what) indicating what assertions are to be listed,
and the reply consists of an error or the requested list.

   The result of requesting a list of rgss3_list_item LABEL objects is a
   list of LFSs supported by the server.  The client can then use the
   LFS list to assert labels via the RPCSEC_GSS_CREATE label assertions.
   See Section 2.7.1.3.

2.8.  Extensibility

   Assertion types may be added in the future by adding arms to the
   "rgss3_assertion_u" union (Section 2.7.1) and the "rgss3_list_item_u"
   union (Section 2.7.2).  Examples of other potential assertion types
   include:

   o  Client-side assertions of identity:

      *  Primary client/user identity.

      *  Supplementary group memberships of the client/user, including
         support for specifying deltas to the membership list as seen on
         the server.

3.  Operational Recommendation for Deployment

   RPCSEC_GSSv3 is a superset of RPCSEC_GSSv2 [RFC5403], which in turn
   is a superset of RPCSEC_GSSv1 [RFC2203], and so can be used in all
   situations where RPCSEC_GSSv2 is used, or where RPCSEC_GSSv1 is used
   and channel-bindings functionality is not needed.  RPCSEC_GSSv3
   should be used when the new functionality is needed.

4.  Security Considerations

   This entire document deals with security issues.

   The RPCSEC_GSSv3 protocol allows for client-side assertions of data
   that is relevant to server-side authorization decisions.  These
   assertions must be evaluated by the server in the context of whether
   the client and/or user are authenticated, whether multi-principal
   authentication was used, whether the client is trusted, what ranges
   of assertions are allowed for the client and the user (separately or
   together), and any relevant server-side policy.

   The security semantics of assertions carried by RPCSEC_GSSv3 are
   application protocol-specific.

   Note that RPCSEC_GSSv3 is not a complete solution for labeling: it
   conveys the labels of actors but not the labels of objects.  RPC
   application protocols may require extending in order to carry object
   label information.

There may be interactions with NFSv4's callback security scheme and
NFSv4.1's [RFC5661] GSS SSV (Secret State Verifier) mechanisms.
Specifically, the NFSv4 callback scheme requires that the server
initiate GSS-API security contexts, which does not work well in
practice; in the context of client-side processes running as the same
user but with different privileges and security labels, the NFSv4
callback security scheme seems particularly unlikely to work well.
NFSv4.1 has the server use an existing, client-initiated RPCSEC_GSS
context handle to protect server-initiated callback RPCs.  The
NFSv4.1 callback security scheme lacks all the problems of the NFSv4
scheme; however, it is important that the server pick an appropriate
RPCSEC_GSS context handle to protect any callbacks.  Specifically, it
is important that the server use RPCSEC_GSS context handles that
authenticate the client to protect any callbacks related to server
state initiated by RPCs protected by RPCSEC_GSSv3 contexts.

As described in Section 2.10.10 of [RFC5661], the client is permitted
to associate multiple RPCSEC_GSS handles with a single SSV GSS
context.  RPCSEC_GSSv3 handles will work well with SSV in that the
man-in-the-middle attacks described in Section 2.10.10 of [RFC5661]
are solved by the new reply verifier (Section 2.3).  Using an
RPCSEC_GSSv3 handle backed by a GSS-SSV mechanism context as a parent
handle in an RPCSEC_GSS_CREATE call, while permitted, is complicated
by the lifetime rules of SSV contexts and their associated RPCSEC_GSS
handles.

5.  IANA Considerations

   This section uses terms that are defined in [RFC5226].

5.1.  New RPC Authentication Status Numbers

   The following new RPC Authentication Status Numbers have been added
   to the IANA registry:

   o  RPCSEC_GSS_INNER_CREDPROBLEM (15)  "No credentials for
      multi-principal assertion inner context user".  See
      Section 2.7.1.1.

   o  RPCSEC_GSS_LABEL_PROBLEM (16)  "Problem with label assertion".
      See Section 2.7.1.3.

   o  RPCSEC_GSS_PRIVILEGE_PROBLEM (17)  "Problem with structured
      privilege assertion".  See Section 2.7.1.4.

   o  RPCSEC_GSS_UNKNOWN_MESSAGE (18)  "Unknown structured privilege
      assertion".  See Section 2.7.1.4.

5.2.  Structured Privilege Name Definitions

   IANA has created a registry called the "RPCSEC_GSS Structured
   Privilege Names Registry".

   Structured privilege assertions (Section 2.7.1.4) are defined by a
   specific RPC application.  The namespace identifiers for these
   assertions (the rp_name) are defined as string names.  The
   RPCSEC_GSSv3 protocol does not define the specific assignment of the
   namespace for these structured privilege assertion names.  The IANA
   registry promotes interoperability where common interests exist.
   While RPC application developers are allowed to define and use
   structured privileges as needed, they are encouraged to register
   structured privilege assertion names with IANA.

   The registry is to be maintained using the Standards Action policy as
   defined in Section 4.1 of [RFC5226].

   Under the RPCSEC_GSS version 3 specification, the name of a
   structured privilege can in theory be up to 2^32 - 1 bytes in length,
   but in practice RPC application clients and servers will be unable to
   handle a string that long.  IANA should reject any assignment request
   with a structured privilege name that exceeds 128 UTF-8 characters.
   To give the IESG the flexibility to set up bases of assignment of
   Experimental Use, the prefix "EXPE" is Reserved.  The structured
   privilege with a zero-length name is Reserved.

   The prefix "PRIV" is allocated for Private Use.  A site that wants to
   make use of unregistered named attributes without risk of conflicting
   with an assignment in IANA's registry should use the prefix "PRIV" in
   all of its structured privilege assertion names.

   Because some RPC application clients and servers have case-
   insensitive semantics, the fifteen additional lower-case and mixed-
   case permutations of each of "EXPE" and "PRIV" are Reserved (e.g.,
   "expe", "expE", and "exPe" are Reserved).  Similarly, IANA must not
   allow two assignments that would conflict if both structured
   privilege names were converted to a common case.

The registry of structured privilege names is a list of assignments, each containing three fields for each assignment.

1.  A US-ASCII string name that is the actual name of the structured privilege.  This name must be unique.  This string name can be 1 to 128 UTF-8 characters long.

2.  A reference to the specification of the RPC-application-defined structured privilege.  The reference can consume up to 256 bytes (or more if IANA permits).

3.  The point of contact of the registrant.  The point of contact can consume up to 256 bytes (or more if IANA permits).

5.2.1.  Initial Registry

The initial registry consists of the three structured privileges defined in [RFC7862].

1.  NAME: copy_to_auth, REFERENCE: RFC 7862, CONTACT: William A.(Andy) Adamson, andros@netapp.com

2.  NAME: copy_from_auth, REFERENCE: RFC 7862, CONTACT: William A.(Andy) Adamson, andros@netapp.com

3.  NAME: copy_confirm_auth, REFERENCE: RFC 7862, CONTACT: William A.(Andy) Adamson, andros@netapp.com

5.2.2.  Updating Registrations

The registrant is always permitted to update the point of contact field.  To make any other change will require Expert Review or IESG Approval.

6.  References

6.1.  Normative References

   [RFC2119]  Bradner, S., "Key words for use in RFCs to Indicate
              Requirement Levels", BCP 14, RFC 2119,
              DOI 10.17487/RFC2119, March 1997,
              <http://www.rfc-editor.org/info/rfc2119>.

   [RFC2203]  Eisler, M., Chiu, A., and L. Ling, "RPCSEC_GSS Protocol
              Specification", RFC 2203, DOI 10.17487/RFC2203,
              September 1997, <http://www.rfc-editor.org/info/rfc2203>.

   [RFC2743]  Linn, J., "Generic Security Service Application Program
              Interface Version 2, Update 1", RFC 2743,
              DOI 10.17487/RFC2743, January 2000,
              <http://www.rfc-editor.org/info/rfc2743>.

   [RFC4506]  Eisler, M., Ed., "XDR: External Data Representation
              Standard", STD 67, RFC 4506, DOI 10.17487/RFC4506,
              May 2006, <http://www.rfc-editor.org/info/rfc4506>.

   [RFC5056]  Williams, N., "On the Use of Channel Bindings to Secure
              Channels", RFC 5056, DOI 10.17487/RFC5056, November 2007,
              <http://www.rfc-editor.org/info/rfc5056>.

   [RFC5403]  Eisler, M., "RPCSEC_GSS Version 2", RFC 5403,
              DOI 10.17487/RFC5403, February 2009,
              <http://www.rfc-editor.org/info/rfc5403>.

   [RFC5661]  Shepler, S., Ed., Eisler, M., Ed., and D. Noveck, Ed.,
              "Network File System (NFS) Version 4 Minor Version 1
              Protocol", RFC 5661, DOI 10.17487/RFC5661, January 2010,
              <http://www.rfc-editor.org/info/rfc5661>.

   [RFC7530]  Haynes, T., Ed., and D. Noveck, Ed., "Network File System
              (NFS) Version 4 Protocol", RFC 7530, DOI 10.17487/RFC7530,
              March 2015, <http://www.rfc-editor.org/info/rfc7530>.

   [RFC7862]  Haynes, T., "Network File System (NFS) Version 4 Minor
              Version 2 Protocol", RFC 7862, DOI 10.17487/RFC7862,
              November 2016, <http://www.rfc-editor.org/info/rfc7862>.

6.2.  Informative References

   [AFS-RXGK]
              Wilkinson, S. and B. Kaduk, "Integrating rxgk with AFS",
              Work in Progress, draft-wilkinson-afs3-rxgk-afs-08,
              May 2015.

   [RFC4949]  Shirey, R., "Internet Security Glossary, Version 2",
              FYI 36, RFC 4949, DOI 10.17487/RFC4949, August 2007,
              <http://www.rfc-editor.org/info/rfc4949>.

   [RFC5226]  Narten, T. and H. Alvestrand, "Guidelines for Writing an
              IANA Considerations Section in RFCs", BCP 26, RFC 5226,
              DOI 10.17487/RFC5226, May 2008,
              <http://www.rfc-editor.org/info/rfc5226>.

   [RFC5531]  Thurlow, R., "RPC: Remote Procedure Call Protocol
              Specification Version 2", RFC 5531, DOI 10.17487/RFC5531,
              May 2009, <http://www.rfc-editor.org/info/rfc5531>.

Authors' Addresses

   William A. (Andy) Adamson
   NetApp
   3629 Wagner Ridge Ct.
   Ann Arbor, MI  48103
   United States of America

   Phone: +1 734 665 1204
   Email: andros@netapp.com


   Nico Williams
   cryptonector.com
   13115 Tamayo Dr.
   Austin, TX  78729
   United States of America

   Email: nico@cryptonector.com