                      NFS Version 4.0 Trunking Update

Abstract

   In NFS version 4.0, the fs_locations attribute informs clients about
   alternate locations of file systems. An NFS version 4.0 client can
   use this information to handle migration and replication of server
   file systems. This document describes how an NFS version 4.0 client
   can also use this information to discover an NFS version 4.0 server's
   trunking capabilities. This document updates RFC 7530.

Status of This Memo

   This is an Internet Standards Track document.

   This document is a product of the Internet Engineering Task Force
   (IETF). It represents the consensus of the IETF community. It has
   received public review and has been approved for publication by the
   Internet Engineering Steering Group (IESG). Further information on
   Internet Standards is available in Section 2 of RFC 7841.

   Information about the current status of this document, any errata,
   and how to provide feedback on it may be obtained at
   https://www.rfc-editor.org/info/rfc8587.

Table of Contents

1. Introduction

   The NFS version 4.0 specification [RFC7530] defines a migration
   feature that enables the transfer of a file system from one server to
   another without disruption of client activity. There were a number
   of issues with the original definition of this feature, now resolved
   with the publication of [RFC7931].

   After a migration event, a client must determine whether state
   recovery is necessary. To do this, it needs to determine whether 1)
   the source and destination server addresses represent the same server
   instance, 2) if the client has already established a lease on the
   destination server for other file systems, and 3) if the destination
   server instance has lock state for the migrated file system.

   As part of addressing this need, [RFC7931] introduces trunking into
   NFS version 4.0 along with a trunking detection mechanism. A
   trunking detection mechanism enables a client to determine whether
   two distinct network addresses are connected to the same NFS version
   4.0 server instance. Without this knowledge, a client unaware of a
   trunking relationship between paths it is using simultaneously is
   likely to become confused in ways described in [RFC7530].

   NFSv4.1 was defined with an integral means of trunking detection,
   which is described in [RFC5661]. NFSv4.0 initially did not have
   trunking detection; it was added by [RFC7931]. Nevertheless, the use
   of the concept of server-trunkability is the same in both protocol
   versions.

   File system migration, replication, and referrals are distinct
   protocol features. However, it is not appropriate to treat each of
   these features in isolation. For example, recovery processing of
   client migration needs to deal with the possibility of multiple
   server addresses in a returned fs_locations attribute. In addition,
   the content of the fs_locations attribute, which provides both
   trunking-related and replication information, may change over
   repeated retrievals, requiring an integrated description of how
   clients are to deal with such changes. The issues discussed in the
   current document relate to the interpretation of the fs_locations
   attribute and to the proper client and server handling of changes in
   fs_locations attribute values.

Therefore, the goals of the current document are as follows:

o To provide NFS version 4.0 with a means of finding addresses that
   are trunkable with a given address, i.e., trunking discovery,
   compatible with the means of trunking detection introduced by
   [RFC7931]. For an explanation of trunking detection and
   discovery, see Section 3.

o To describe how NFS version 4.0 clients are to handle the presence
   of multiple network addresses associated with the same server when
   recovering from a replication and migration event.

o To describe how NFS version 4.0 clients are to handle changes in
   the contents of returned fs_locations attributes, including those
   that indicate changes in the responding NFS version 4.0 server's
   trunking configuration.

The current document pursues these goals by presenting a set of
updates to [RFC7530], as summarized in Sections 5 and 6.

2. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
"SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and
"OPTIONAL" in this document are to be interpreted as described in
BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all
capitals, as shown here.

3. Terminology

Most of the terms related to handling the fs_locations attribute are
appropriately defined in Section 5.1. However, there are a few
terminological issues regarding the use of terms outside the context
of text updating [RFC7530] that are explained in this section. Note
that the definitions of trunking-related terms in Section 5.1 apply
throughout this document, including in explanatory sections that will
not replace any text in [RFC7530].

Regarding network addresses and the handling of trunking, we use the
following terminology:

o Each NFSv4 server is assumed to have a set of IP addresses to
   which NFSv4 requests may be sent by clients. These are referred
   to as the server's "network addresses". Access to a specific
   server network address might involve the use of multiple network
   ports, since the ports to be used for particular types of
   connections might be required to be different.

   o Clients may establish connections to NFSv4 servers via one of
      several connection types, supporting the NFSv4 protocol layered on
      top of an RPC stream transport, as described in [RFC5531], or on
      top of RPC-over-RDMA, as described in [RFC8166]. The combination
      of a server network address and a particular connection type is
      referred to as a "server endpoint".

   o Each network address, when combined with a pathname providing the
      location of a file system root directory relative to the
      associated server root filehandle, defines a file system network
      access path.

   o Two network addresses connected to the same server are said to be
      server-trunkable. Unlike subsequent NFSv4 minor versions, NFSv4.0
      recognizes only a single type of trunking relationship between
      addresses.

   Discussion of the term "replica" is complicated for a number of
   reasons. Even though the term is used in explaining the issues in
   [RFC7530] that need to be addressed in the current document, a full
   explanation of this term requires explanation of related terms
   connected to the fs_locations attribute, which is provided in
   Section 5.1 of the current document.

   The term is also used in previous documents about NFSv4.0 (i.e.,
   [RFC7530] and [RFC7931]) with a meaning different from that in the
   current document. In these documents, each replica is identified by
   a single network access path. However, in the current document, a
   set of network access paths that have server-trunkable network
   addresses and the same root-relative file system pathname is
   considered to be a single replica with multiple network access paths.
   Although [RFC7931] enables an NFSv4.0 client to determine whether two
   network addresses are server-trunkable, it never describes the
   addresses as connected to a single replica, in effect leaving the
   approach established in [RFC7530].

   Note that this document, except when explaining problems in
   [RFC7530], always uses the new definition, including in text intended
   to replace existing sections of [RFC7530].

4. Document Organization

   The sections of the current document are divided into four types
   based on how they relate to the eventual updating of the NFS version
   4.0 specification. Once this update is published, NFS version 4.0
   will be specified by multiple documents that need to be read together
   until such time as a consolidated replacement specification is
   produced.

   o The base specification [RFC7530]

   o The migration-related update [RFC7931]

   o This document [RFC8587]

   The section types are as follows. See Appendix A for a
   classification of each section of the current document.

   o An explanatory section does not contain any material that is meant
      to update the specification of NFS version 4.0. Such sections may
      contain an explanation about why and how changes are to be made,
      but they do not include any text that is to update [RFC7530] or
      appear in an eventual consolidated document.

   o A replacement section contains text that is to replace and thus
      supersede text within [RFC7530] and then appear in an eventual
      consolidated document. The titles of the replacement sections
      indicate what section of [RFC7530] is to be replaced.

   o An additional section contains text that, although not replacing
      anything in [RFC7530], will be part of the specification of NFS
      version 4.0 and will be expected to be part of an eventual
      consolidated document. The titles of the additional sections
      provide an indication of where the new section would appear when
      consolidated with [RFC7530].

   o An editing section contains some text that replaces text within
      [RFC7530], although the entire section will not consist of such
      text and will include other text as well. Such sections make
      relatively minor adjustments in the existing NFS version 4.0
      specification, which are expected to be reflected in an eventual
      consolidated document. Generally, such replacement text appears
      as a quotation, possibly taking the form of an indented set of
      paragraphs.

Additional and replacement sections sometimes contain references to
the "current document" by which RFC 8587 is meant. When those
sections are incorporated in a consolidated document, those
references will need to be updated to refer to the appropriate
sections in that new document.

5. Changes within Section 8 of RFC 7530

Most of the updates to [RFC7530] that provide support for trunking
using the fs_locations attribute apply to Section 8 ("Multi-Server
Namespace") of that document. In the following list, the replacing
section refers to its numbering in this document.

o Section 5.1 replaces Section 8.1 ("Location Attributes") of
   [RFC7530]. The text in the original section has been reorganized
   and extended to explicitly allow the use of fs_locations to
   provide trunking-related information that appropriately interacts
   with the migration, replication, and referral features of
   fs_locations. Terminology used to describe the interactions is
   added.

o Section 5.2 updates Section 8.4 ("Uses of Location Information")
   of [RFC7530]. This section comprises the bulk of the updates.
   Each paragraph of Section 8.4 and its subsections have been
   reviewed to clarify the provision of trunking-related information
   using the fs_locations attribute.

   * Section 5.2.1 replaces the introductory material within
      Section 8.4 of [RFC7530], i.e., the material within Section 8.4
      exclusive of subsections.

   * Section 5.2.2 is to be added as a new subsection of Section 8.4
      before the updated Section 8.4.1 of [RFC7530]. In a
      consolidated document, it would appear as Section 8.4.1.

   * Section 5.2.3 is to be added as a new subsection of Section 8.4
      before the updated Section 8.4.1 of [RFC7530]. In a
      consolidated document, it would appear as Section 8.4.2.

   * Section 5.2.4 replaces Section 8.4.1 ("File System
      Replication") of [RFC7530]. In a consolidated document, it
      would appear as Section 8.4.3.

   * Section 5.2.5 replaces Section 8.4.2 ("File System Migration")
      of [RFC7530]. In a consolidated document, it would appear as
      Section 8.4.4.

      * Section 5.2.6 is to be added as a new subsection of Section 8.4
         before Section 8.4.3 of [RFC7530]. In a consolidated document,
         it would appear as Section 8.4.5, while the existing
         Section 8.3 would appear as Section 8.4.6.

   o Section 5.3 replaces Section 8.5 ("Location Entries and Server
      Identity") of [RFC7530]. The last paragraph of the existing
      section has been removed.

5.1. Updated Section "Location Attributes" (Currently Section 8.1)

   The fs_locations attribute allows specification of file system
   locations where the data corresponding to a given file system may be
   accessed. This attribute represents such file system instances as a
   server address target (as either a DNS hostname representing one or
   more network addresses or as a single literal network address)
   together with the path of that file system within the associated
   single-server namespace. Individual fs_locations entries can express
   trunkable addresses, locations of file system replicas on other
   servers, migration targets, or pure referrals.

   We introduce the following terminology:

   o "Trunking" is a situation in which multiple network addresses are
      connected to the same NFS server. Network addresses connected to
      the same NFS server instance are said to be "server-trunkable".

   o "Trunking detection" refers to ways of confirming that two
      distinct network addresses are connected to the same NFSv4 server
      instance.

   o Trunking discovery is a process by which a client using one
      network address can obtain other candidate addresses that are
      server-trunkable with it.

   Regarding terminology relating to GETATTR attributes used in trunking
   discovery and other multi-server namespace features:

   o Location attributes include only the fs_locations GETATTR
      attribute.

   o Location entries (fs_location4, defined in [RFC7530],
      Section 2.2.6) are the individual file system locations in the
      fs_locations attribute (defined in [RFC7530], Section 2.2.7). A
      file system location entry designates a set of network addresses
      to which clients may establish connections. The entry may
      designate multiple such addresses because the server hostname may
      map to multiple network addresses and because multiple connection

types may be used to communicate with each specified network
address. Such addresses provide multiple ways of connecting to a
single server.

Clients use the NFSv4.0 trunking detection mechanism [RFC7931] to
confirm that such addresses are connected to the same server. The
client can ignore non-confirmed trunking relationships and treat
the corresponding addresses as connected to different servers.

o File system location elements are derived from file system
   location entries. If a file system location entry specifies a
   network address, there is only a single corresponding location
   element. When a file system location entry contains a hostname,
   the client resolves the hostname, producing one file system
   location element for each of the resulting network addresses.
   Issues regarding the trustworthiness of hostname resolutions are
   further discussed in Section 7 of the current document.

o All file system location elements consist of a file system
   location address, which is the network address of an interface to
   a server, and an fs_name, which is the location of the file system
   within the server's pseudo-fs.

o If the server has no pseudo-fs and only has a single exported file
   system at the root filehandle, the fs_name may be empty.

5.2. Updates to "Uses of Location Information" (Currently Section 8.4)

   The subsections below provide replacement sections for existing
   sections within Section 8.4 of [RFC7530] or new subsections to be
   added to that section.

5.2.1. Updates to the Introductory Text of the Current Section 8.4

   Together with the possibility of absent file systems, the
   fs_locations attribute bears file system locations and a number of
   important facilities that enable reliable, manageable, and scalable
   data access.

   When a file system is present on the queried server, this attribute
   can provide a set of alternate locations that clients may use to
   access the file system, when necessary. Provision of such alternate
   file system locations is referred to as "replication" and is further
   described in Section 5.2.4 of the current document.

   When alternative file system locations are provided, they may
   represent distinct physical copies of the same file system data or
   separate NFS server instances that provide access to the same

physical file system. Another possible use of the provision of
multiple file system location entries is trunking, wherein the file
system location entries do not, in fact, represent different servers
but rather are distinct network paths to the same server.

A client may use file system location elements simultaneously to
provide higher-performance access to the target file system. This
can be done using trunking, although the use of multiple replicas
simultaneously is possible. To enable simultaneous access, the
client utilizes trunking detection and/or discovery, further
described in Section 5.2.2 of the current document, to determine a
set of network paths that are server-trunkable with the path
currently being used to access the file system. Once this
determination is made, requests may be routed across multiple paths
using the existing state management mechanism.

Multiple replicas may also be used simultaneously, typically when
accessing read-only datasets. In this case, each replica requires
its own state management. The client performs multiple file opens to
read the same file content from multiple replicas.

When a file system is present and subsequently becomes absent,
clients can be given the opportunity to have continued access to
their data at an alternative file system location. Transfer of the
file system contents to the new file system location is referred to
as "migration". The client's responsibilities in dealing with this
transition depend on the specific nature of the new access path as
well as how and whether data was, in fact, migrated. See Sections
5.2.5 and 5.2.6 of the current document for details.

The fs_locations attribute can designate one or more remote file
system locations in place of an absent file system. This is known as
a "referral". A particularly important case is that of a "pure
referral", in which the absent file system has never been present on
the NFS server. Such a referral is a means by which a file system
located on one server can redirect clients to file systems located on
other servers, thus enabling the creation of a multi-server
namespace.

Because client support for the fs_locations attribute is OPTIONAL, a
server may (but is not required to) take action to hide migration and
referral events from such clients by acting as a proxy, for example.

5.2.2. New Subsection Titled "Trunking Discovery and Detection"
        (Becomes Section 8.4.1)

   "Trunking" is a situation in which multiple distinct network
   addresses are associated with the same NFS server instance. As a
   matter of convenience, we say that two network addresses connected to
   the same NFS server instance are server-trunkable. Section 5.4 of
   [RFC7931] explains why NFSv4 clients need to be aware of the NFS
   server identity to manage lease and lock states effectively when
   multiple connections to the same server exist.

   "Trunking detection" refers to a way for an NFSv4 client to confirm
   that two independently acquired network addresses are connected to
   the same NFSv4 server. Section 5.8 of [RFC7931] describes an
   OPTIONAL means by which it can be determined whether two network
   addresses correspond to the same NFSv4.0 server instance. Without
   trunking detection, an NFSv4.0 client has no other way to confirm
   that two network addresses are server-trunkable.

   In the particular context of NFS version 4.0, trunking detection
   requires that the client support the uniform client ID string (UCS)
   approach, described in Section 5.6 of [RFC7931]. Any NFSv4.0 client
   that supports migration or trunking detection needs to present a
   uniform client ID string to all NFSv4.0 servers. If it does not do
   so, it will be unable to perform trunking detection.

   "Trunking discovery" is the process by which an NFSv4 client, using a
   hostname or one of an NFSv4 server's network addresses, can obtain
   other candidate network addresses that are trunkable with the NFSv4
   server's network address, i.e., a set of addresses that might be
   connected to the same NFSv4 server instance. An NFSv4.0 client can
   discover server-trunkable network addresses in a number of ways:

   o An NFS server's hostname is provided either at mount time or in a
      returned file system location entry. A DNS query of this hostname
      can return more than one network address. The returned network
      addresses are candidates for trunking.

   o Location entries returned in an fs_locations attribute can specify
      network addresses. These network addresses are candidates for
      trunking.

   When there is a means of trunking detection available, an NFSv4.0
   client can confirm that a set of network addresses corresponds to the
   same NFSv4.0 server instance; thus, any of them can be used to access
   that server.

5.2.3. New Subsection Titled "Location Attributes and Selection of
        Connection Type" (Becomes Section 8.4.2)

   NFS version 4.0 may be implemented using a number of different types
   of connections:

      Stream connections may be used to provide RPC service, as
      described in [RFC5531].

      RDMA-capable connections may be used to provide RPC service, as
      described in [RFC8166].

   Because of the need to support multiple connection types, clients
   face the issue of determining the proper connection type to use when
   establishing a connection to a server network address. The
   fs_locations attribute provides no information to support selection
   of the connection type. As a result, clients supporting multiple
   connection types need to attempt to establish a connection on various
   connection types, allowing it to determine, via a trial-and-error
   approach, which connection types are supported.

   If a client strongly prefers one connection type, it can perform
   these attempts serially in order of declining preference. Once there
   is a successful attempt, the established connection can be used.
   Note that with this approach, network partitions can result in a
   sequence of long waits for a successful connection.

   To avoid waiting when there is at least one viable network path
   available, simultaneous attempts to establish multiple connection
   types are possible. Once a viable connection is established, the
   client discards less-preferred connections.

5.2.4. Updated Section "File System Replication" (Becomes Section 8.4.3
        Retitled "File System Replication and Trunking"

   On first access to a file system, the client should obtain the value
   of the set of alternative file system locations by interrogating the
   fs_locations attribute. Trunking discovery and/or detection can then
   be applied to the file system location entries to separate the
   candidate server-trunkable addresses from the replica addresses that
   provide alternative locations of the file system. Server-trunkable
   addresses may be used simultaneously to provide higher performance
   through the exploitation of multiple paths between the client and
   target file system.

   In the event that server failure, communication problems, or other
   difficulties make continued access to the current file system
   impossible or otherwise impractical, the client can use the

alternative file system locations as a way to maintain continued
access to the file system. See Section 5.2.6 of the current document
for more detail.

5.2.5. Updated Section "File System Migration" (Becomes Section 8.4.4)

When a file system is present and becomes absent, clients can be
given the opportunity to have continued access to their data at an
alternative file system location specified by the fs_locations
attribute. Typically, a client will be accessing the file system in
question, get an NFS4ERR_MOVED error, and then use the fs_locations
attribute to determine the new location of the data. See
Section 5.2.6 of the current document for more detail.

Such migration can help provide load balancing or general resource
reallocation. The protocol does not specify how the file system will
be moved between servers. It is anticipated that a number of
different server-to-server transfer mechanisms might be used, with
the choice left to the server implementer. The NFSv4 protocol
specifies the method used to communicate the migration event between
the client and server.

When the client receives indication of a migration event via an
NFS4ERR_MOVED error, data propagation to the destination server must
have already occurred. Once the client proceeds to access the
alternate file system location, it must see the same data. Where
file systems are writable, a change made on the original file system
must be visible on all migration targets. Where a file system is not
writable but represents a read-only copy (possibly periodically
updated) of a writable file system, similar requirements apply to the
propagation of updates. Any change visible in the original file
system must already be effected on all migration targets to avoid any
possibility that a client, in effecting a transition to the migration
target, will see any reversion in the file system state.

5.2.6. New Subsection Titled "Interaction of Trunking, Migration, and
        Replication" (Becomes Section 8.4.5)

   When the set of network addresses on a server changes in a way that
   would affect a file system location attribute, there are several
   possible outcomes for clients currently accessing that file system.
   NFS4ERR_MOVED is returned only when the server cannot satisfy a
   request from the client, whether because the file system has been
   migrated to a different server or is only accessible at a different
   trunked address on the same server, or for some other reason. In
   cases 1 and 2 below, NFS4ERR_MOVED is not returned.

   1. When the list of network addresses is a superset of that
       previously in effect, there is no need for migration or any other
       sort of client adjustment. Nevertheless, the client is free to
       use an additional address in the replacement list if that address
       provides another path to the same server. Alternatively, the
       client may treat that address as it does a replica -- to be used
       if the current server addresses become unavailable.

   2. When the list of network addresses is a subset of that previously
       in effect, immediate action is not needed if an address missing
       in the replacement list is not currently in use by the client.
       The client should avoid using that address to access that file
       system in the future, whether the address is for a replica or an
       additional path to the server being used.

   3. When an address being removed is one of a number of paths to the
       current server, the client may continue to use it until
       NFS4ERR_MOVED is received. This is not considered a migration
       event unless the last available path to the server has become
       unusable.

   When migration does occur, multiple addresses may be in use on the
   server prior to migration, and multiple addresses may be available
   for use on the destination server.

   With regard to the server in use, a return of NFS4ERR_MOVED may
   indicate that a particular network address is no longer to be used,
   without implying that migration of the file system to a different
   server is needed. Clients should not conclude that migration has
   occurred until confirming that all network addresses known to be
   associated with that server are not usable.

It should be noted that the need to defer this determination is not
absolute. If a client is not aware of all network addresses for any
reason, it may conclude that migration has occurred when it has not
and treat a switch to a different server address as if it were a
migration event. This is harmless since the use of the same server
via a new address will appear as a successful instance of transparent
state migration.

Although significant harm cannot arise from this misapprehension, it
can give rise to disconcerting situations. For example, if a lock
has been revoked during the address shift, it will appear to the
client as if the lock has been lost during migration. When such a
lock is lost, it is the responsibility of the destination server to
provide for its recovery via the use of an fs-specific grace period.

With regard to the destination server, it is desirable for the client
to be aware of all valid network addresses that can be used to access
the destination server. However, there is no need for this to be
done immediately. Implementations can process the additional file
system location elements in parallel with normal use of the first
valid file system location entry found to access the destination.

Because a file system location attribute may include entries relating
to the current server, the migration destination, and possible
replicas to use, scanning for available network addresses that might
be trunkable with addresses the client has already seen could
potentially be a long process. To keep this process as short as
possible, servers that provide information about trunkable network
paths are REQUIRED to place file system location entries that
represent addresses usable with the current server or a migration
target before those associated with replicas.

This ordering allows a client to cease scanning for trunkable file
system location entries once it encounters a file system location
element whose fs_name differs from the current fs_name or whose
address is not server-trunkable with the address it is currently
using. Although the possibility exists that a client might
prematurely cease scanning for trunkable addresses when receiving a
location attribute from an older server that does not follow the
ordering constraint above, the harm is expected to be limited since
such servers would not be expected to present information about
trunkable server access paths.

5.3. Updated Section "Location Entries and Server Identity"
      (Section 8.5)

   As mentioned above, a single file system location entry may have a
   server address target in the form of a DNS hostname that resolves to
   multiple network addresses; it is also possible for multiple file
   system location entries to have their own server address targets that
   reference the same server.

   When server-trunkable addresses for a server exist, the client may
   assume that for each file system in the namespace of a given server
   network address, file systems at corresponding namespace locations
   exist for each of the other server-trunkable network addresses. It
   may do this even in the absence of explicit listing in fs_locations.
   Such corresponding file system locations can be used as alternative
   locations, just as those explicitly specified via the fs_locations
   attribute.

   If a single file system location entry designates multiple server IP
   addresses, the client should choose a single one to use. When two
   server addresses are designated by a single file system location
   entry and they correspond to different servers, this normally
   indicates some sort of misconfiguration. The client should avoid
   using such file system location entries when alternatives are
   available. When they are not, the client should pick one of the IP
   addresses and use it without using others that are not directed to
   the same server.

6. Updates to RFC 7530 outside Section 8

   Since the existing description of NFS4ERR_MOVED in Section 13.1.2.4
   of [RFC7530] does not take proper account of trunking, it needs to be
   modified by replacing the first two sentences of the description with
   the following material:

      The file system that contains the current filehandle object cannot
      be accessed using the current network address. It may be
      accessible using other network addresses connected to the same
      server, it may have been relocated to another server, or it may
      never have been present.

7. Updates to the Security Considerations Section of RFC 7530

   The Security Considerations section of [RFC7530] needs the additions
   below to properly address some aspects of trunking discovery,
   referral, migration, and replication.

The possibility that requests to determine the set of network
addresses corresponding to a given server might be interfered with
or have their responses corrupted needs to be taken into account.

o When DNS is used to convert NFS server hostnames to network
   addresses and DNSSEC [RFC4033] is not available, the validity
   of the network addresses returned cannot be relied upon.
   However, when the client uses RPCSEC_GSS [RFC7861] to access
   NFS servers, it is possible for mutual authentication to detect
   invalid server addresses. Other forms of transport layer
   security (e.g., [RFC8446]) can also offer strong authentication
   of NFS servers.

o Fetching file system location information SHOULD be performed
   using RPCSEC_GSS with integrity protection, as previously
   explained in the Security Considerations section of [RFC7530].
   Making a request of this sort without using strong integrity
   protection permits corruption during the transit of returned
   file system location information. The client implementer needs
   to recognize that using such information to access an NFS
   server without use of RPCSEC_GSS (e.g., by using AUTH_SYS as
   defined in [RFC5531]) can result in the client interacting with
   an unverified network address that is posing as an NFSv4
   server.

o Despite the fact that [RFC7530] REQUIRES "implementations" to
   provide "support" for the use of RPCSEC_GSS, it cannot be
   assumed that use of RPCSEC_GSS is always possible between any
   particular client-server pair.

o Returning only network addresses to a client that has no
   trusted DNS resolution service can hamper its ability to use
   RPCSEC_GSS.

Therefore, an NFSv4 server SHOULD present file system location
entries that correspond to file systems on other servers using
only hostnames. This enables the client to interrogate the
fs_locations on the destination server to obtain trunking
information (as well as replica information) using RPCSEC_GSS with
integrity, validating the hostname provided while ensuring that
the response has not been corrupted.

When RPCSEC_GSS is not available on an NFS server, returned file
system location information is subject to corruption during
transit and cannot be relied upon. In the case of a client being
directed to another server after NFS4ERR_MOVED, this could vitiate
the authentication provided by the use of RPCSEC_GSS on the
destination. Even when RPCSEC_GSS authentication is available on

the destination, this server might validly represent itself as the
server to which the client was erroneously directed. Without a
way to decide whether the server is a valid one, the client can
only determine, using RPCSEC_GSS, that the server corresponds to
the hostname provided, with no basis for trusting that server.
The client should not use such unverified file system location
entries as a basis for migration, even though RPCSEC_GSS might be
available on the destination server.

When a file system location attribute is fetched upon connecting
with an NFSv4 server, it SHOULD, as stated above, be done using
RPCSEC_GSS with integrity protection.

When file system location information cannot be protected in
transit, the client can subject it to additional filtering to
prevent the client from being inappropriately directed. For
example, if a range of network addresses can be determined that
ensure that the servers and clients using AUTH_SYS are subject to
appropriate constraints (such as physical network isolation and
the use of administrative controls within the operating systems),
then network addresses in this range can be used, with others
discarded or restricted in their use of AUTH_SYS.

When neither integrity protection nor filtering is possible, it is
best for the client to ignore trunking and replica information or
simply not fetch the file system location information for these
purposes.

To summarize considerations regarding the use of RPCSEC_GSS in
fetching file system location information, consider the following
recommendations for requests to interrogate location information,
with interrogation approaches on the referring and destination
servers arrived at separately:

o The use of RPCSEC_GSS with integrity protection is RECOMMENDED
   in all cases, since the absence of integrity protection exposes
   the client to the possibility of the results being modified in
   transit.

o The use of requests issued without RPCSEC_GSS (e.g., using
   AUTH_SYS), while undesirable, might be unavoidable in some
   cases. Where the use of returned file system location
   information cannot be avoided, it should be subject to
   filtering to eliminate untrusted network addresses. The
   specifics will vary depending on the degree of network
   isolation and whether the request is to the referring or
   destination servers.

Privacy considerations relating to uniform client strings (UCS)
versus non-uniform client strings (non-UCS), discussed in
Section 5.6 of [RFC7931], are also applicable to their usage for
trunking detection in NFS version 4.0.

8. IANA Considerations

This document has no IANA actions.

9. Updates to the References Section in RFC 7530

The following references should be added to the Normative References
section of [RFC7530]:

[RFC7931] Noveck, D., Ed., Shivam, P., Lever, C., and B. Baker,
           "NFSv4.0 Migration: Specification Update", RFC 7931,
           DOI 10.17487/RFC7931, July 2016,
           <https://www.rfc-editor.org/info/rfc7931>.

[RFC8166] Lever, C., Ed., Simpson, W., and T. Talpey, "Remote
           Direct Memory Access Transport for Remote Procedure
           Call Version 1", RFC 8166, DOI 10.17487/RFC8166,
           June 2017, <https://www.rfc-editor.org/info/rfc8166>.

The following references should be added to the Informative
References section of [RFC7530]:

[RFC4033] Arends, R., Austein, R., Larson, M., Massey, D.,
           and S. Rose, "DNS Security Introduction and
           Requirements", RFC 4033, DOI 10.17487/RFC4033,
           March 2005, <https://www.rfc-editor.org/info/rfc4033>.

[RFC7861] Adamson, A. and N. Williams, "Remote Procedure Call
           (RPC) Security Version 3", RFC 7861, DOI 10.17487/RFC7861,
           November 2016, <https://www.rfc-editor.org/info/rfc7861>.

[RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol
           Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018,
           <https://www.rfc-editor.org/info/rfc8446>.

10. References

10.1. Normative References

   [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate
              Requirement Levels", BCP 14, RFC 2119,
              DOI 10.17487/RFC2119, March 1997,
              <https://www.rfc-editor.org/info/rfc2119>.

   [RFC5531] Thurlow, R., "RPC: Remote Procedure Call Protocol
              Specification Version 2", RFC 5531, DOI 10.17487/RFC5531,
              May 2009, <https://www.rfc-editor.org/info/rfc5531>.

   [RFC7530] Haynes, T., Ed. and D. Noveck, Ed., "Network File System
              (NFS) Version 4 Protocol", RFC 7530, DOI 10.17487/RFC7530,
              March 2015, <https://www.rfc-editor.org/info/rfc7530>.

   [RFC7931] Noveck, D., Ed., Shivam, P., Lever, C., and B. Baker,
              "NFSv4.0 Migration: Specification Update", RFC 7931,
              DOI 10.17487/RFC7931, July 2016,
              <https://www.rfc-editor.org/info/rfc7931>.

   [RFC8166] Lever, C., Ed., Simpson, W., and T. Talpey, "Remote Direct
              Memory Access Transport for Remote Procedure Call Version
              1", RFC 8166, DOI 10.17487/RFC8166, June 2017,
              <https://www.rfc-editor.org/info/rfc8166>.

   [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC
              2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174,
              May 2017, <https://www.rfc-editor.org/info/rfc8174>.

10.2. Informative References

   [RFC4033] Arends, R., Austein, R., Larson, M., Massey, D., and S.
              Rose, "DNS Security Introduction and Requirements",
              RFC 4033, DOI 10.17487/RFC4033, March 2005,
              <https://www.rfc-editor.org/info/rfc4033>.

   [RFC5661] Shepler, S., Ed., Eisler, M., Ed., and D. Noveck, Ed.,
              "Network File System (NFS) Version 4 Minor Version 1
              Protocol", RFC 5661, DOI 10.17487/RFC5661, January 2010,
              <https://www.rfc-editor.org/info/rfc5661>.

   [RFC7861] Adamson, A. and N. Williams, "Remote Procedure Call (RPC)
              Security Version 3", RFC 7861, DOI 10.17487/RFC7861,
              November 2016, <https://www.rfc-editor.org/info/rfc7861>.

   [RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol
              Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018,
              <https://www.rfc-editor.org/info/rfc8446>.

Appendix A. Section Classification

   All sections of the current document are considered explanatory with
   the following exceptions.

   o Sections 5.1, 5.2.4, 5.2.5, and 5.3 are replacement sections.

   o Sections 5.2.2, 5.2.3, and 5.2.6 are additional sections.

   o Sections 5.2.1, 6, 7, and Section 9 are editing sections.

Authors' Addresses

   Charles Lever (editor)
   Oracle Corporation
   United States of America

   Email: chuck.lever@oracle.com

   David Noveck
   NetApp
   United States of America

   Email: davenoveck@gmail.com